

Outsmarting Time:

Foundation Models for Zero-Shot

Forecasting



Tom Altenborg & Kristian Ertresvåg

Supervisor: Jonas Moss

BI Norwegian Business School – MSc Business Analytics

Acknowledgements

We would like to express our gratitude to everyone who has supported us throughout our master's thesis, which marks the completion of our master's program in Business Analytics at BI Norwegian Business School.

First and foremost, we thank our supervisor, Jonas Moss, for the helpful guidance and feedback. The expertise and encouragement were crucial for steering us in the right direction early on.

We also appreciate the support from our loved ones, who have encouraged us along the way. Your unwavering support and belief have been a constant source of motivation.

We would also like to thank Didier Lopes, founder and CEO of OpenBB, for inspiring our thesis.

Finally, we extend our thanks to colleagues and peers for camaraderie, discussions and insights. Their contributions and shared experiences have enriched our time as students and made the process more enjoyable.

Abstract

This thesis evaluates the performance of three foundation models – Chronos, TimeGPT and Moirai – in the context of the M3-Competition. The models are based on the Transformer architecture, a wildly successful architecture in the field of Artificial Intelligence for both Large Language Models and Computer Vision. We evaluate the foundation models in zero-shot forecasting, as the models are not fitted to the data before inference, against the original competitors in the competition.

Our evaluation shows promise for using foundation models for time series forecasting. Overall, Chronos outperforms the two other foundation models and places fifth in our ranking of the 23 models for all data. The foundation models display better results for monthly data than the other time frequencies in the dataset, with Chronos being the top performer of all evaluated models for this subset. Our findings suggest that foundation models can be a viable approach to forecasting, and we recommend considering Chronos as the preferred model for forecasting monthly data.

All code and data used in this thesis are available in this [GitHub repository](#).

Contents

1	Introduction	6
2	Literature Review	7
3	Theoretical Framework	9
3.1	Time Series	9
3.2	Time Series Forecasting	10
3.2.1	Naïve Forecasting	10
3.2.2	Exponential Smoothing	10
3.2.3	ARIMA Models	14
3.2.4	ARARMA Models	17
3.2.5	Decomposition Models	17
3.2.6	Expert Models	20
3.3	Deep Learning	20
3.3.1	Neural Networks	20
3.3.2	Recurrent Neural Networks	23
3.3.3	Long Short-Term Memory Networks	25
3.3.4	Convolutional Neural Networks	26
3.4	Transformer-Based Models	27
3.4.1	Architecture	27
3.4.2	Inference	31
3.4.3	Transfer Learning	32
3.4.4	Transformers for Time Series Forecasting	33
3.5	TimeGPT	34
3.5.1	Model Architecture	35
3.5.2	Training	36
3.5.3	Approach to Time Series Forecasting	36
3.5.4	Available Methods and Capabilities	37
3.6	Chronos	37
3.6.1	Model Architecture	37
3.6.2	Training	38
3.6.3	Approach to Time Series Forecasting	39
3.6.4	Available Methods and Capabilities	40

3.7	MOIRAI	40
3.7.1	Model Architecture	40
3.7.2	Training	41
3.7.3	Approach to Time Series Forecasting	42
3.7.4	Available Methods and Capabilities	42
3.8	Model Evaluation Metrics	42
3.8.1	sMAPE	43
3.8.2	MdAPE	44
3.8.3	MdRAE	45
3.8.4	RMSE	45
3.8.5	MASE	46
4	Methodology	47
4.1	Introduction	47
4.2	Competition Selection	47
4.2.1	The M3 Competition	48
4.3	Evaluation	50
4.3.1	Evaluation Metrics	50
4.3.2	Benchmark Model	51
4.3.3	Leaderboard	51
4.3.4	Time Frequencies	52
4.3.5	Categories	52
4.3.6	Comparison Against Benchmark	52
4.4	Statistical Tests	53
4.5	Implementation - TimeGPT	53
4.5.1	Access	53
4.5.2	Reproducibility	54
4.5.3	API Data Integration	54
4.6	Implementation - Chronos & Moirai	54
4.6.1	Data Preparation	54
4.6.2	Model Configuration	55
4.6.3	Forecast Generation and Storage	55
5	Data	56

5.1	Data Collection	56
5.2	Data Quality	56
5.3	Data Description	56
5.4	Data Preparation	58
5.5	Data Limitations	58
6	Results	59
6.1	Competition Results	59
6.1.1	Leaderboard Position	60
6.2	Evaluation	60
6.2.1	Time Frequencies	61
6.2.2	Categories	64
6.2.3	Benchmark	65
6.3	Robustness of Results	66
7	Discussion	69
7.1	Discussion of Results	69
7.2	Limitations	70
7.3	Future Research	71
8	Conclusion	73
	Bibliography	81
	Appendix A	82
	Appendix B	84
	Appendix C	85
	Appendix D	86
	Appendix E	88

1 Introduction

The field of time series forecasting has seen significant development with the introduction of advanced machine learning models. Transformer-based foundation models, in particular, have drawn considerable interest for their ability to handle complex data and scale efficiently with data. This thesis examines the performance of three such models—*Chronos*, *TimeGPT*, and *Moirai*—within the M3-Competition, a competition for zero-shot time series forecasting.

Time series forecasting is crucial in various domains, including finance, economics, supply chain management, and meteorology. Accurate forecasts can lead to better decision-making, optimized operations, and improved strategic planning. Traditionally, ARIMA and exponential smoothing models have been central to time series analysis. However, the rise of machine learning has introduced alternatives that promise improved accuracy and adaptability. Initially developed for natural language processing tasks, transformers have shown remarkable potential in handling sequential data, including time series. Their ability to model long-range dependencies without the limitations of recurrent architectures such as Long Short-Term Memory networks positions them as strong candidates for time series forecasting. This research aims to evaluate the effectiveness of foundation models in a competitive forecasting environment.

This thesis focuses on evaluating the performances of *Chronos*, *TimeGPT*, and *Moirai*. The M3-Competition serves as an excellent platform for this evaluation, offering a varied collection of time series data across different frequencies and categories to thoroughly benchmark model performance. Additionally, the study explores how these models respond to variations in data frequency and category, highlighting their strengths and weaknesses. This evaluation also covers aspects such as ease of access, reproducibility, and practical integration of the models, ensuring the results are relevant for real-world applications.

In summary, this thesis aims to enrich the knowledge of foundation models in time series forecasting. By evaluating their performance against traditional benchmarks and assessing their adaptability across different data contexts, we seek to offer insights to both researchers and practitioners. Our findings will underscore both the potential and the areas for improvement of these models.

2 Literature Review

Time series forecasting has undergone significant technological advancements and theoretical developments since the mid-20th century. From the initial statistical methods to today's advanced deep learning models, the field has grown alongside computational capabilities and developments in analytical techniques. From the 1950s, Exponential Smoothing methods originated from the work of Brown (1959); Brown (1963), Holt (1957) and Winters (1960). While these methods became popular among practitioners, they struggled to attract attention from statisticians due to the absence of a solid statistical foundation. It was Muth (1960) who first placed exponential smoothing into a statistical context, a path later pursued by the likes of Box and Jenkins (1970). The techniques received further exposure and credibility through reviews and studies by Gardner and McKenzie (1985), who extensively evaluated the methods developed until then. This demonstrated that Simple Exponential Smoothing could be viewed as a model based on a single error source, reinforcing the statistical basis of exponential smoothing (De Gooijer & Hyndman, 2006).

Yule (1927) introduced the concept that every time series could be considered the realization of a stochastic process, meaning the data points can be assumed to be outcomes from a probabilistic model. This idea led to the first formulations of autoregressive (AR) and moving average (MA) models. The integration of AR with MA models into autoregressive moving average (ARMA) models was further innovated by Slutsky (1937) and N. and Wold (1939). Box and Jenkins (1970) provided further work on these concepts and introduced the Box-Jenkins methodology, which optimizes the application of ARMA models and introduced autoregressive integrated moving average (ARIMA) models for handling non-stationary data. The Box-Jenkins approach enormously impacted the theory and practice of time series analysis and forecasting in modern times (De Gooijer & Hyndman, 2006). Hyndman and Athanasopoulos (2021), and earlier editions, provides a thorough overview of these traditional methods and applications. With the innovations of computers, ARIMA models became more practical for practitioners, but as access to computational power has increased, machine learning models have also become a popular choice. Boosting algorithms were introduced to improve the accuracy of predictive models. Freund and Schapire (1997) developed AdaBoost as a boosting technique for classification problems. This laid the groundwork for Friedman (2001), which introduced gradient boosting. Gradient boosting later became the foundations for Chen and Guestrin (2016) with XGBoost and Ke et al. (2017) with LightGBM, two of the most popular machine learning models for forecasting in recent years.

Krizhevsky et al. (2012) is regarded as the breakthrough that showed the power of deep learning by

achieving remarkable results in the ImageNet Large Scale Visual Recognition Challenge. The authors advanced the performance of image classification tasks by using a Convolutional Neural Network (CNN) architecture named AlexNet. Although this specific model does not apply to time series, and related work such as Lecun et al. (1998) is focused on image processing and recognition, it has popularized CNNs, which have been important for using deep learning in time series forecasting. An example of this is Lai et al. (2018), which presents the Long and Short-term Time-series network, a model that combines both CNNs and Recurrent Neural Networks (RNNs) for time series forecasting. Today's most popular AI models for natural language processing (NLP), as Jiang et al. (2023) with Mistral 7B, are based on the transformer architecture introduced by Vaswani et al. (2017). Earlier work as Hochreiter and Schmidhuber (1997) on Long Short-Term Memory and Bai et al. (2018) on Temporal Convolutional Networks has also contributed to the foundation for utilizing deep learning in time series forecasting.

The early contributions of Brown and Yule to today's sophisticated deep learning models mark a tremendous evolution in the complexity of modeling techniques in time series forecasting. However, a meta-analysis by Green and Armstrong (2015) indicates that increased model complexity does not guarantee superior performance over simpler methods. This observation finds support in the success of the Theta model, developed by Assimakopoulos and Nikolopoulos (2000), which demonstrated competitive results in the M3 competition against much more elaborate models, as presented by Makridakis and Hibon (2000). This underscores the value of benchmarking new forecasting models against simpler alternatives, providing a critical perspective on their relative performance. While the landscape of forecasting and artificial intelligence is ever-evolving, the future might change these dynamics. Nonetheless, the current state of research suggests a cautious approach to assuming complexity as a hallmark of effectiveness in forecasting models.

3 Theoretical Framework

3.1 Time Series

A *time series* is a collection of data points recorded in sequential order over time. It can be *continuous*, recording data continuously over time, or *discrete*, capturing data at specific time intervals. The series type—continuous or discrete—refers to the nature of the time axis, not the variable measured, which can be either discrete or continuous in both types. Discrete time series emerge in three ways: sampling from a continuous series, aggregating over time periods, or inherently discrete occurrences. Chatfield (2000), pp. 11–12 notes that these data points are typically recorded at regular time intervals. A sampled continuous series may be exemplified by hourly temperature recordings, where variables are sequentially logged from an intrinsically continuous series. An example of aggregated data could be the cumulative count of COVID-19 cases since the start of the pandemic, where the data is compiled from continuous records to the total number of infections.

Time series analysis aims to describe data through summary statistics and visual methods, identify appropriate statistical models for the data-generating process, forecast future values, and inform decision-making. Chatfield (2000), pp. 12–13 explains that time series data is distinctive because of temporal dependencies, meaning successive observations often depend on each other, necessitating analysis that accounts for this order. These temporal dependencies are also evident in text data, as word order matters.

Classical time series analysis decomposes data variations into four components: *seasonal variation* (annual patterns), *trend* (consistent upward or downward movement), *other cyclic variations* (non-annual cycles), and *irregular fluctuations* (unpredictable changes) (Chatfield, 2000, pp. 13–14). A popular example of seasonal variation is the increase in ice cream sales during the summer months. In contrast, other cyclic variations, like economic recessions, typically extend over multiple years of variable lengths. The global population's continuous growth over many decades is an example of a trend. Irregular fluctuations, which are often labelled as "noise", are not part of a systematic pattern.

Time series that lack seasonal variation and a trend is referred to as *stationary*. With stationary time series, the time we observe the data does not matter for the statistical properties. When working with such data, one could use a sample from any given time to forecast the next period. When seasonality or a trend is evident, the data is not stationary, as the time of the observations will matter for the following observation. However, this does not include non-annual cycles since the cycles do not have fixed lengths. Hyndman and Athanasopoulos (2021) clarify that the time of the observations does not

provide any information about the following observation, as cycles can peak and end at any time.

3.2 Time Series Forecasting

In recent years, the field of time series forecasting has witnessed substantial advancements, evolving in methodology and application. To provide a focused and relevant overview of the theoretical framework, our examination will primarily concentrate on the most relevant theoretical foundations for the statistical forecasting models used in the M3-Competition. All models discussed in this section are capable of zero-shot time series forecasting, meaning they can be used without exogenous variables.

3.2.1 Naïve Forecasting

Forecasting techniques do not necessarily need to be complicated. A method that has worked surprisingly well in time series forecasting, is the *naïve method*:

$$\hat{y}_{(T+h|T)} = y_T \quad (1)$$

With this method, the forecast $\hat{y}_{(T+h|T)}$ is simply equal to the last observation in the series y_T . Hyndman and Athanasopoulos (2021) describe this as a random walk forecast, which is optimal when the movements are unpredictable and equally likely to go up or down.

An alternative to the naïve method is the *seasonal naïve method*, which can be applied when working with highly seasonal data. The forecast for this method will be equal to the last observation from the same season:

$$\hat{y}_{(T+h|T)} = y_{T+h-m(k+1)} \quad (2)$$

Here, $y_{T+h-m(k+1)}$ is the last observation from the same season as $\hat{y}_{(T+h|T)}$. With m being the seasonal period and k is the number of prior periods. $T+h-m(k+1)$ calculates the seasonal index for the last observation from the same season as $T+h$ (Hyndman & Athanasopoulos, 2021).

3.2.2 Exponential Smoothing

Exponential smoothing, a technique tracing its roots to the mid-20th century, has been a cornerstone in time series forecasting. Initially regarded as a set of heuristic methods, exponential smoothing has come quite a way from then on to acquire a robust statistical basis. This approach is characterized by

its ability to weigh past observations with an exponentially decreasing importance, making it effective in smoothing data and predicting short-term trends in a time series. The introduction of state space models to exponential smoothing marked a significant advancement, offering a more systematic and flexible framework for dealing with various types of data, including those with complex patterns and seasonality. De Gooijer and Hyndman (2006) note that this enhancement improved the method's forecasting accuracy and expanded its applicability across diverse fields, ranging from economics to supply chain management.

Simple Exponential Smoothing (SES) represents the most basic form of exponential smoothing and is suitable for datasets lacking evident trends or seasonal fluctuations. The formula for SES can be given as:

$$\hat{y}_{T+1} = \alpha y_T + (1 - \alpha) \hat{y}_T \quad (3)$$

This equation predicts the value at the next step, $t + 1$, in a time series, based on known values up to and including t . The forecast is denoted as $\hat{y}_{(t+1|t)}$ which is calculated by adjusting the previous forecast $\hat{y}_{(t|t-1)}$ with the most recent observation y_t , using a weighting factor α . The factor α , called the smoothing constant, ranges between 0 and 1 and balances the influence of the recent observation against the prior forecast. If α is closer to 1, the model places more emphasis on the latest observed data, making the forecast more sensitive to recent changes. Conversely, if α is closer to 0, the model gives more weight to the historical forecasts, thus smoothing out random fluctuations and highlighting long-term trends. The term $(1 - \alpha)$ scales the previous forecast's influence, inversely related to α (Hyndman & Athanasopoulos, 2021).

Alternatively, SES can be represented in component form. This form provides no additional value but is helpful as we add more components. The level, ℓ_t , is introduced in the component form and is the flat forecast for SES, where h is the number of time periods forecasted.

$$\hat{y}_{(t+h|t)} = \ell_t \quad (4)$$

$$\ell_t = \alpha y_t + (1 - \alpha) \ell_{t-1} \quad (5)$$

Holt's Linear Trend Method Building on SES, Charles Holt (1957) introduced *Holt's linear trend method*, also called Double Exponential Smoothing (DES), as an extension that factors in trends in the data (Hyndman & Athanasopoulos, 2021). This method has three equations:

$$\hat{y}_{(t+h|t)} = \ell_t + hb_t \quad (6)$$

$$\ell_t = \alpha y_t + (1 - \alpha)(\ell_{t-1} + b_{t-1}) \quad (7)$$

$$b_t = \beta^*(\ell_t - \ell_{t-1}) + (1 - \beta^*)b_{t-1} \quad (8)$$

In DES, ℓ_t is similar to (5) but is now also adjusted for the previous trend b_{t-1} . The trend b_t is estimated based on a change in the level, $\ell_t - \ell_{t-1}$, the previous trend b_{t-1} and a smoothing factor β^* . The smoothing factor β^* has the same characteristics as α but is unique for the trend as it is possible to differentiate the two based on preference or calculations, as we will come back to later. The forecast (6) is now a linear function of h with the level as intercept and trend as slope (Hyndman & Athanasopoulos, 2021).

As mentioned, Holt's linear method is a linear function, meaning the trend will be constant. This can cause overshooting for longer forecast horizons. Gardner and McKenzie (1985) modified the method to *dampen* the trend:

$$\hat{y}_{(t+h|t)} = \ell_t + \sum_{i=1}^h \phi^i b_t \quad (9)$$

$$\ell_t = \alpha y_t + (1 - \alpha)(\ell_{t-1} + \phi b_{t-1}) \quad (10)$$

$$b_t = \beta^*(\ell_t - \ell_{t-1}) + (1 - \beta^*)\phi b_{t-1} \quad (11)$$

This *dampening trend method* introduces an autoregressive-dampening parameter ϕ , dampening the trend as long as $0 < \phi < 1$. If the parameter ϕ is close to 1, the forecast will approach the forecast of Holt's linear method. Conversely, if the parameter ϕ is 0, the trend will be eliminated from the forecast. This parameter could also be higher than 1 to create an exponential trend. Gardner and McKenzie (1985) notes this as an autoregressive-damping (AD) parameter. It is the equivalent of the autoregressive term in an ARIMA model, which we will come back to. Hyndman and Athanasopoulos (2021) observe that in practice, the parameter ϕ typically ranges between 0.8 and 0.98, as the method is intended to dampen trends. However, they note that the dampening strongly affects smaller values, making a ϕ smaller than 0.8 rarely desirable.

Holt-Winters Seasonal Method Holt's linear trend method is further extended to capture seasonality in *Holt-Winters seasonal method* by Holt (1957) and Winters (1960). This method has two variations, whereas the additive method is preferred for approximately constant seasonal variations.

Holt-Winters' additive method is given as:

$$\hat{y}_{(t+h|t)} = \ell_t + hb_t + s_{t+h-m(k-1)} \quad (12)$$

$$\ell_t = (\alpha(y_t) - s_{t-m}) + (1 - \alpha)(\ell_{t-1} + b_{t-1}) \quad (13)$$

$$b_t = \beta^*(\ell_t - \ell_{t-1}) + (1 - \beta^*)b_{t-1} \quad (14)$$

$$s_t = \gamma(y_t - \ell_{t-1} - b_{t-1}) + (1 - \gamma)s_{t-m} \quad (15)$$

With this method, $s_{t+h-m(k-1)}$ is added to the forecast function (9) and is the seasonal index adjusted for the length of the season cycle, with m being the seasons within one year, and the forecast horizon h . The seasonal component s_t in (12) is again similar to the earlier equation with a smoothing factor γ . However, it is calculated with respect to s_{t-m} to capture the last season instead of using the last observation. The smoothing factor γ also differs from the previously mentioned α and β^* as it usually is restricted to $0 \leq \gamma \leq 1 - \alpha$ (Hyndman & Athanasopoulos, 2021).

Hyndman and Athanasopoulos (2021) recommend the multiplicative variation of Holt-Winters' method when the seasonality changes proportional to the time series' level. The multiplicate variation is:

$$\hat{y}_{(t+h|t)} = (\ell_t + hb_t)s_{t+h-m(k-1)} \quad (16)$$

$$\ell_t = \alpha \frac{y_t}{s_{t-m}} + (1 - \alpha)(\ell_{t-1} + b_{t-1}) \quad (17)$$

$$b_t = \beta^*(\ell_t - \ell_{t-1}) + (1 - \beta^*)b_{t-1} \quad (18)$$

$$s_t = \gamma \left(\frac{y_t}{\ell_{t-1} - b_{t-1}} \right) + (1 - \gamma)s_{t-m} \quad (19)$$

This variation will return a percentage as the seasonal component, and the forecast will be calculated by multiplying this percentage with the linear function $\ell_t + hb_t$, which we recognize as (6). This variation captures the relative changes in seasons to calculate the forecast instead of the additive method, where seasonality is added on in absolute numbers.

As with Holt's linear method, it is possible to dampen the trend for the Holt-Winters seasonal method for both the additive and multiplicative variation. With multiplicate seasonality, the model is:

$$\hat{y}_{(t+h|t)} = \left[\ell_t + \sum_{i=1}^h \phi^i b_t \right] s_{t+h-m(k-1)} \quad (20)$$

$$\ell_t = \alpha \frac{y_t}{s_{t-m}} + (1 - \alpha) (\ell_{t-1} + \phi b_{t-1}) \quad (21)$$

$$b_t = \beta^* (\ell_t - \ell_{t-1}) + (1 - \beta^*) \phi b_{t-1} \quad (22)$$

$$s_t = \gamma \left(\frac{y_t}{\ell_{t-1} - \phi b_{t-1}} \right) + (1 - \gamma) s_{t-m} \quad (23)$$

Here, h is exchanged with $\sum_{i=1}^h \phi^i$ to multiply the parameter for dampening the trend ϕ to the forecast. This works as the dampening trend method for Holt's linear method. Further equations for the level, trend and seasonality are modified with the parameter ϕ multiplied with b_{t-1} to dampen the trend for b_{t-1} in each part of the model (Hyndman & Athanasopoulos, 2021).

Optimization The application of each of the exponential smoothing methods requires the forecaster to choose the smoothing parameters and initial values. Hyndman and Athanasopoulos (2021) assert that while these values can be picked based on experience or intuition, estimating the values based on data is more reliable. With mathematical programming, it is possible to decide both smoothing parameters and initial values with a calculation of the sum of the squared residuals (SSE):

$$SSE = \sum_{t=1}^T (y_t - \hat{y}_{(t|t-1)})^2 = \sum_{t=1}^T e_t^2 \quad (24)$$

This function calculates the sum of the squared residuals where the residuals are $e_t = y_t - \hat{y}_{(t|t-1)}$. With this approach, the SSE would be the objective function to be minimized subject to the relevant equations given by the method and the smoothing parameters and initial values as decision variables.

3.2.3 ARIMA Models

Next off is the popular *Auto-Regressive Integrated Moving Average* (ARIMA) models. ARIMA models are particularly renowned for their versatility, capable of modeling a wide range of time series with varying levels of complexity. The Box-Jenkins methodology encompasses model identification, parameter estimation, and diagnostic checking and provides a systematic framework for ARIMA modeling, ensuring robustness and accuracy in forecasts. This methodology has been influential in econometrics and has also found applications across diverse domains such as environmental science and inventory control. De Gooijer and Hyndman (2006) highlight that ARIMA models continue to evolve and remain relevant, showing how dynamic forecasting methods are and how crucial they are for understanding patterns over time. Central to ARIMA are the three concepts of integrating *autoregressive* (AR) and *moving average* (MA) components, along with *differencing* (I) to achieve stationarity, thus allowing for the effective modeling of both short-term and long-term dependencies

in data. The “Integrated” in ARIMA refers to the process of differencing.

Autoregression models forecast the variable of interest using a linear combination of past observations. It is similar to multiple regression models, where the variable of interest is forecasted using a linear combination of predictors but with lagged values of the variable of interest as predictors. With the order p , an Autoregressive model, referred to as $AR(p)$, can be written as:

$$y_t = c + \varepsilon_t + \sum_{i=1}^p \phi_i y_{t-i} \quad (25)$$

Hyndman and Athanasopoulos (2021) explain that in this model, c is the intercept, while ϕ_i are the parameters that reflect the influence of past observations on the current value. The error term ε_t represents white noise, accounting for fluctuations not explained by past values.

In contrast, the Moving Average model uses past forecast errors in a regression-like model. It is not an actual regression since we do not observe the errors, but it is similar:

$$y_t = c + \varepsilon_t + \sum_{j=1}^q \theta_j \varepsilon_{t-j} \quad (26)$$

In (26), ε_t is again the white noise, and c is the intercept, but q is the order. The parameter θ_j represents the impact of the past error terms on the current value. This is referred to as an $MA(q)$ model.

For both Autoregressive and Moving Average models, the given parameters can be restricted to certain constraints. In its simplest form, for $AR(1)$ and $MA(1)$, both with only one value for the parameter, the parameter can be restricted to only be between -1 and 1. However, increasing the order of the models complicates the constraints. Hyndman and Athanasopoulos (2021) point out that the white noise, meaning a sequence of random variables unexplained by the parameters with a mean of zero and a constant variance, for the models, is assumed. The variance of the error term for both models will influence the scale of the series, not patterns.

The $ARMA(p, q)$ model combines these two models into one model to use both last observations and past forecast errors to predict the current value. The model can be given as:

$$y_t = c + \varepsilon_t + \sum_{i=1}^p \phi_i y_{t-i} + \sum_{j=1}^q \theta_j \varepsilon_{t-j} \quad (27)$$

Here, the current value y_t is calculated based on the lagged values of y_t and the lagged errors of ε_t . Again, ϕ_i and θ_j holds the influence of these values as (25) and (26) respectively.

The process of differencing is the last component of the ARIMA model. Differencing is a method to make non-stationary time series stationary by computing the differences between consecutive observations, removing the level and stabilizing the mean. Differencing can be shown as:

$$y'_t = y_t - y_{t-1} \quad (28)$$

Here, the difference y'_t is simply calculated by subtracting the previous observation y_{t-1} from the given observation y_t . Alternatively, Hyndman and Athanasopoulos (2021) demonstrates that it is possible to difference the data based on seasons instead of the last observation. To achieve this, exchange y_{t-1} with y_{t-m} in (28), where m is the number of seasons, to calculate the difference y'_t between the last observation with the last observation from the same season. This is called *seasonal differencing*.

However, this may not always be enough to obtain a stationary time series, which can be handled by differencing the data a second time, called *second-order differencing*:

$$y_t^* = y'_t - y'_{t-1} = (y_t - y_{t-1}) - (y_{t-1} - y_{t-2}) \quad (29)$$

With second-order differencing, we obtain y_t^* which is the difference in differences between y'_t and y'_{t-1} in which we recognize the definition of from (28).

We can now combine the three concepts of autoregression, differencing and moving averages into the non-seasonal ARIMA model. The full model is given as:

$$y'_t = c + \varepsilon_t + \sum_{i=1}^p \phi_i y'_{t-i} + \sum_{j=1}^q \theta_j \varepsilon_{t-j} \quad (30)$$

This is the $ARIMA(p, d, q)$ model where the three predictors are the order of the autoregression (p), the degree of first differencing (d) and the order of the moving average (q). This is similar to (27) but with differencing. Hyndman and Athanasopoulos (2021) note that modern programming libraries can automatically select values for the predictors using the ARIMA function, simplifying the deployment of these models. However, they also caution that because predictors significantly influence the forecast's behavior, it may be necessary to use additional procedures such as evaluating autocorrelations or applying domain knowledge to choose appropriate values.

$ARIMA(p, d, q)$ can be extended further to the seasonal ARIMA model, denoted as $ARIMA(p, d, q) (P, D, Q)_m$

where P, D, Q is the same predictors p, d and q but the values is specifically for the seasonal part of the ARIMA model and m is again the number of periods within a season (Hyndman & Athanasopoulos, 2021). However, this falls out of the scope of this thesis as it is not used in the M3 Competition.

3.2.4 ARARMA Models

Parzen (1982) introduced the Auto-Regressive Auto-Regressive Moving average (ARARMA). This method uses a transformed series defined by an autoregressive operator instead of differencing.

$$\tilde{y}_t = c + \varepsilon_t + \sum_{i=1}^p \phi_i \tilde{y}_{t-i} + \sum_{j=1}^q \theta_j \varepsilon_{t-j} \quad (31)$$

The current value \tilde{y}_t is calculated based on the constant c , the error term ε_t and the AR and MA components of predictors p and q with coefficients ϕ_i and θ_j , respectively. The MA component uses error terms as earlier, but the AR components now calculate based on a time series \tilde{y}_t . The time series \tilde{y}_t is generated through an autoregressive process as (25):

$$\tilde{y}_t = c + \varepsilon_t + \sum_{i=1}^p \phi_i y_{t-i} \quad (32)$$

The ARARMA model is similar to ARIMA models, employing the $ARMA(p, q)$ model for calculating the current value. What sets ARARMA apart is its autoregressive approach to transform the series for modeling purposes, rather than using the raw or differenced series as is in ARMA or ARIMA models. This distinction allows the ARARMA method to model the time series directly through its own past values, with an alternative approach to achieve stationarity.

3.2.5 Decomposition Models

A fundamental approach in forecasting and time series analysis is *time series decomposition*. The primary goal of decomposition methods is to identify and separate a time series into three fundamental components: trend, seasonality, and remainder (or residuals). When we assume a multiplicative decomposition for a time series, we can write:

$$y_t = S_t * T_t * R_t \quad (33)$$

In this equation, y_t is our time series where it is the product of a seasonal component S_t , a trend-cycle component T_t and a remainder component R_t at period t . Hyndman and Athanasopoulos (2021)

expresses that a time series can be conceptualized as a product of different components, enabling a deeper examination of each component. There are several decomposition methods, and they can also be additive:

$$y_t = S_t + T_t + R_t \quad (34)$$

Here, the time series y_t is decomposed into the sum of the same three factors. An additive approach assumes that the components add up to the observed data. This is suitable for data when the factors are roughly constant over time t , while a multiplicative approach is more suitable when the variations change proportionally to the level.

Classical Decomposition Method Before presenting the Theta-model, it is important to introduce the classical decomposition method in its multiplicative form. It is a relatively simple method and is often used as a starting point for other decomposition methods. Hyndman and Athanasopoulos (2021) put forth the four steps of the method: (1) compute the trend-cycle component T_t , (2) detrend the series, (3) estimate seasonal component S_t and (4) calculate the remainder component R_t .

In the first step, the trend-cycle component T_t is computed with a moving average of order m , also called m -MA. This can be written as:

$$\hat{T}_t = \frac{1}{m} \sum_{j=-k}^k y_{t+j} \quad (35)$$

In this equation, $m = 2k + 1$. This is similar to the formula for calculating a mean, but it is specified to calculate the average for the time series within k periods of t where m is the seasonal period. This means that, for example, with quarterly data ($m = 4$), this equation will calculate the average for the centered moving average over one year, which, in this case, consists of four quarters. The result is centered as it takes the average across a symmetric time window around each data point, rather than leading or lagging values for point of interest y_t . However, this result will not be centered if m is even. This is because the calculation would straddle two central periods rather than align with one. To handle this, a second moving average is computed, referred to as $2 \times m - MA$, which averages the values from the initial m -MA to center the moving average on actual time periods. This additional step shifts the focus of the moving average to the midpoint of the original m periods when m is even. Hyndman and Athanasopoulos (2021) highlight a drawback of this approach: it cannot estimate \hat{T}_t at both the beginning and the end of the time series because centered moving averages require data on

both sides of the point to be smoothed.

The second step of the method is to use the trend-cycle component to detrend the series. The detrended time series is obtained by $\frac{y_t}{\hat{T}_t}$. The third step is to use the detrended series to calculate the average of detrended values for each season, which gives \hat{S}_t . This allows for the method's final step, which is to calculate the remainder component. As we already have the other components, we can solve for $\hat{R}_t = \frac{y_t}{(\hat{T}_t \hat{S}_t)}$. Again, as this approach cannot obtain an estimate for \hat{T}_t at both the beginning and at the end of the time series, the detrended series and therefore \hat{S}_t and \hat{R}_t cannot be estimated accurately either for the same points in time.

Theta-Model The Theta-model, introduced by Assimakopoulos and Nikolopoulos (2000), proposes a different decomposition approach as it decomposes seasonally adjusted series into short- and long-term components. Central to the model is the Theta-coefficient θ , which is applied to a time series subject to second order differencing y_t^* to obtain a Theta-line L_t :

$$L_t = \theta y_t^* \quad (36)$$

Using this model, different values for θ serve different purposes. Setting θ equal to 0 will transform the series to a linear regression line, revealing the overall trend of the series. Having $\theta > 1$ will amplify the short-term behavior of the data while $0 < \theta < 1$ will return a line for the data with a dampened effect. It is then possible to decompose the initial time series into two or more Theta-lines to combine short- and long-term fluctuations into a product. Assimakopoulos and Nikolopoulos (2000) demonstrate this with an example that decomposes the time series into two Theta-lines, one where $\theta = 0$ and another where $\theta = 2$. This will have one line for the linear trend and one for short-term behavior. The lines can be further extrapolated by other methods, like simple exponential smoothing, although this is only relevant for $L_t (\theta = 2)$ in this example. With these two lines, we can calculate:

$$M_t = \frac{1}{2} (L_t (\theta = 0) + L_t (\theta = 2)) \quad (37)$$

This will return the resulting model M_t , which is the average of the two different Theta-lines $L_t (\theta = 0)$ and $L_t (\theta = 2)$, allowing the forecaster to obtain a model based both on linear trend and short-term fluctuations. Assimakopoulos and Nikolopoulos (2000) clarifies that this is only one of several possibilities with this approach, as it is possible to add more lines, change the weighting, extrapolate with

different methods, and use different Theta-lines for different forecasting horizons to obtain a model that fits the dataset.

3.2.6 Expert Models

Expert models, or expert systems, will in this thesis refer to both commercial forecasting software and automated, or semi-automated, rule-based forecasting techniques. In this thesis, the expert models will have in common that they chose a method for each time series based on the characteristics of the data. These methods do not necessarily fall within the scope of the traditional techniques discussed. This categorization is consistent with that of the M3-Competition.

3.3 Deep Learning

Deep learning saw a rise in popularity in 2012 when Krizhevsky et al. (2012) introduced AlexNet, which made a breakthrough on the ImageNet dataset, surpassing the nearest competitor by approximately 10% better accuracy. Larger datasets, better computation, and bolder model designs helped deep neural networks, now known as deep learning, gain momentum in research and applicability. These networks learn directly from raw data, and through training, higher-level features emerge in a process called representation learning (Bommasani et al., 2022).

3.3.1 Neural Networks

Artificial neural networks are based on concepts on how the human brain works, where early studies go as far back as the 1940s, with McCulloch and Pitts (1943) exploration on how neurons could work. The building block of deep neural networks begins with the *Multi-Layer Perceptron* (MLP) model, a fully connected *feedforward network*.

The architecture consists of nodes, or neurons, organized into three main layers: (1) the input layer, which receives the data, (2) one or more hidden layers, where data transformation and computation occur, and (3) the output layer, which produces the results. Each neuron in a layer is connected to every neuron in the subsequent layer through parameters known as weights, with additional parameters called biases contributing to the neurons' output.

The first step in an MLP is to initialize the weights and biases, typically setting them to small random values. This randomness is important to break symmetry and ensure the model learns effectively. There are various strategies like *Xavier* (Glorot & Bengio, 2010) or *He* initialization (He et al., 2015), depending on the linearity or nonlinearity of the activation function, that can improve accuracy and

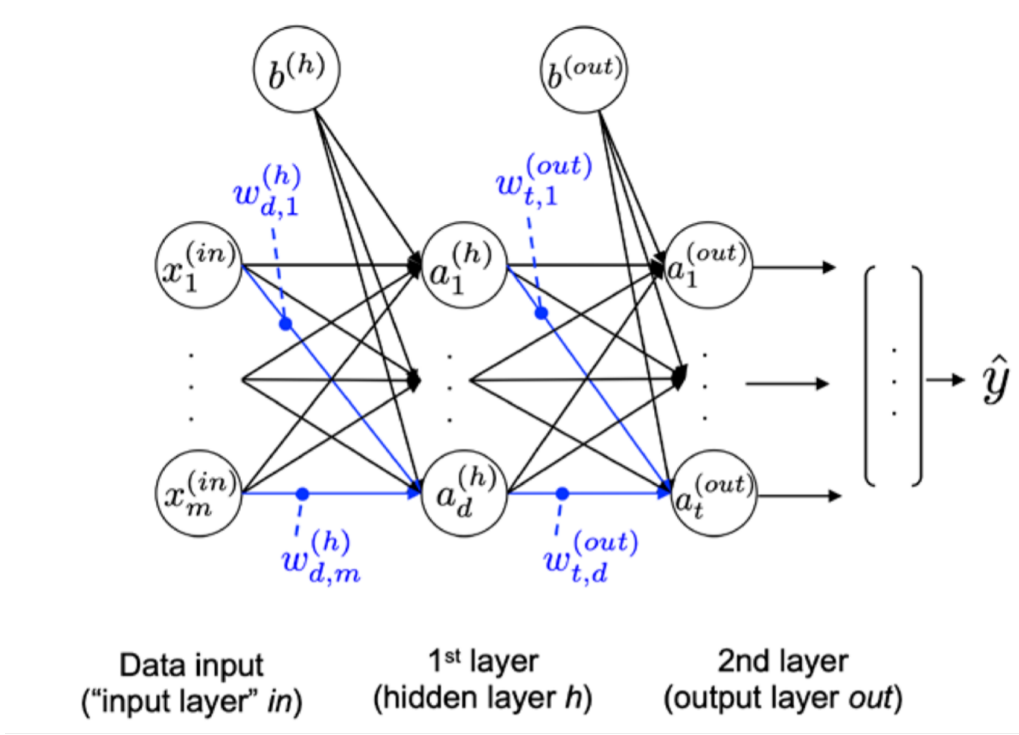


Figure 1: Illustration of a Multilayer feedforward Neural Network (Raschka et al., 2022)

training time.

The next step is forward propagation, which involves passing the input data through the network from the input layer, through the hidden layers, to the output layer. At each layer, the input is transformed by a weighted sum followed by an *activation function*. This process results in the network's prediction based on the current state of its weights. Activation functions introduce nonlinearities into the model, enabling it to learn complex patterns. Commonly used activation functions include:

- **ReLU:** $\sigma(z) = \max(0, z)$
- **GELU:** $\sigma(z) = z * P(Z \leq z) = z \phi(z) = z * \frac{1}{2} \left[1 + \operatorname{erf}\left(\frac{z}{\sqrt{2}}\right) \right]$ if $Z \sim N(0, 1)$

where z is the net input from the network. Gaussian Error Linear Units (GELU) introduced by Hendrycks and Gimpel (2016), is presently the most common due to its performance. However, Raschka et al. (2022) highlights that the choice of activation function depends on whether it is a classification or a regression problem, network architecture, training speed and network performance.

After forward propagation, the network's performance is evaluated using a loss function (sometimes referred to as cost function or objective function). The choice of the loss function is an important hyperparameter of a neural network, where some, such as Dräger and Dunkelau (2022), argue it is the most important. Some commonly used loss functions for regression problems include *Mean Squared Error* and *Mean Absolute Error*:

- **Mean Squared Error (MSE):** $MSE = \frac{1}{n} \sum_{i=1}^n \left(y^{(i)} - \hat{y}^{(i)} \right)^2$
- **Mean Absolute Error (MAE):** $MAE = \frac{1}{n} \sum_{i=1}^n \left| y^{(i)} - \hat{y}^{(i)} \right|$

Each loss function has its own pros and cons, including sensitivity to outliers, challenges in convergence, and more. For instance, the MAE is more resistant to outliers than the MSE because it calculates the average of absolute differences between predicted and actual values. This avoids the exaggerated impact that squaring the errors in MSE has on large deviations. Additionally, it is desirable for a loss function to be differentiable, as gradient descent relies on this property to update weights and biases effectively during backpropagation, steering the model toward minimizing the loss.

Backpropagation is the process of calculating the gradient of the loss function with respect to each weight and bias in the network. This is done by applying the chain rule of calculus and propagating the error backwards through the network. If we modify the MSE slightly with a normalization factor for convenience, as shown by Raschka et al. (2022), the loss function for a simple classification problem can be defined as:

$$L(\mathbf{w}, b) = \frac{1}{2n} \sum_{i=1}^n \left(y^{(i)} - \sigma(z^{(i)}) \right)^2 \quad (38)$$

where \mathbf{w} is the weight vector, b is the bias term, $y^{(i)}$ is the actual class and $\sigma(z^{(i)})$ is the activation function applied to the net input. The net input has the following definition:

$$z^{(i)} = \mathbf{w}^T \mathbf{x}^{(i)} + b = \sum_{j=1}^m \mathbf{w}_j * \mathbf{x}_j^{(i)} + b \quad (39)$$

where $\mathbf{x}^{(i)}$ is the feature vector of the i -th sample and m is the total number of features. This can then be used to compute the partial derivative of this loss function with respect to each weight w_j in the weight vector and bias:

$$\frac{\partial L}{\partial w_j} = -\frac{1}{n} \sum_{i=1}^n \left(y^{(i)} - \sigma(z^{(i)}) \right) \sigma'(z^{(i)}) x_j^{(i)} \quad (40)$$

$$\frac{\partial L}{\partial b} = -\frac{1}{n} \sum_{i=1}^n \left(y^{(i)} - \sigma(z^{(i)}) \right) \sigma'(z^{(i)}) \quad (41)$$

Here, $\sigma'(z^{(i)})$ is the partial derivative of the activation function. The process repeats numerous times to train the network and minimize the errors. It works as the derivative tells what direction the function is growing, and we take a step in the opposite direction, where the learning rate and the slope of the gradient determine the size. Thus, the updating of the weight and bias terms are defined as:

$$\Delta w_j = -\eta \frac{\partial L}{\partial w_j} \quad (42)$$

$$\Delta b = -\eta \frac{\partial L}{\partial b} \quad (43)$$

where η is the learning rate. There are also more optimized gradient descent methods, such as mini-batch and stochastic gradient descent, designed to speed up the process. Furthermore, several issues should be considered when using gradient descent, such as the learning rate η . If too small, the optimization can get “stuck” in local minima; if too large, it can overshoot the global minima. Stochastic gradient descent is often used in practice due to its performance and computational desirability.

Adam optimisation has emerged as particularly popular among the various optimization techniques that build upon traditional gradient descent. Adam optimization, introduced by Kingma and Ba (2014), is a technique for gradient descent optimization that has gained popularity in deep learning due to its computational efficiency, low memory requirements, and minimal manual tuning. It does so by combining two previous ideas, Momentum and RMSprop. Momentum accelerates gradient descent by considering previous gradients to smooth the update. At the same time, RMSprop modifies the learning rate for each parameter by dividing it by an exponentially decaying average of squared gradients. Adam also incorporates a decaying learning rate mechanism, gradually reducing the learning rate as training progresses. These features make Adam particularly effective for training deep neural networks.

3.3.2 Recurrent Neural Networks

Feed-forward neural networks excel at mapping one-to-one relationships between inputs and outputs. However, they struggle with sequential data, where the output depends on a sequence of previous elements. *Recurrent Neural Networks* (RNN) possess an internal memory state, enabling them to retain information from previous inputs. This design allows them to process sequences, making them suitable for tasks such as language translation and time series forecasting. In this context, we will concentrate on what is known as the *Vanilla RNN*, which was once commonly employed for

language-related tasks.

At the heart of RNNs is the internal memory state at each time step. Staudemeyer and Morris (2019) explains that this mechanism is made possible with hidden layers with a self-looping workflow, meaning the hidden layers can remember previous inputs and use these to make predictions. Predictions are made with the current input, the learned weights and biases, and the stored memory of the hidden layers.

Staudemeyer and Morris (2019) underline that RNNs need different training than normal feedforward networks, where the most common method is backpropagation through time (BPTT). Werbos (1990) introduced BPTT, however, mainly within the task of pattern recognition and other methods than neural networks. Raschka et al. (2022) describe BPTT as distinct from standard backpropagation because it feeds the output from one time step back into the network, allowing errors to be recalculated and temporal dependencies to be accounted for. It can do so by computing the gradients for each step while considering the weights of the loss at each step. Using the simplified definition above for the loss function (38) with time indices, the loss function is:

$$L(\mathbf{w}, b) = \frac{1}{2n} \sum_{i=1}^n \sum_{t=1}^T \left(y_t^{(i)} - \sigma(z_t^{(i)}) \right)^2 \quad (44)$$

Where $y_t^{(i)}$ is the actual output at time t and $\sigma(z_t^{(i)})$ is the predicted at time t . As it must consider output from previous layers as well, the gradient will be defined as:

$$\frac{\partial L^{(t)}}{\partial \mathbf{w}_{hh}} = \sum_{k=1}^t \left(\frac{\partial L^{(t)}}{\partial \mathbf{o}^{(t)}} * \frac{\partial \mathbf{o}^{(t)}}{\partial h^{(t)}} * \frac{\partial h^{(t)}}{\partial h^{(k)}} * \frac{\partial h^{(k)}}{\partial \mathbf{w}_{hh}} \right) \quad (45)$$

$L^{(t)}$ represents the loss at time t , $\mathbf{o}^{(t)}$ is the output of the network at time t and $h^{(t)}$ is hidden state of the network at time t . \mathbf{w}_{hh} is a matrix that connects the hidden states with the hidden layers at the next time step, sharing the weights across the network, thereby making it recurrent. This process fine-tunes the model's ability to learn from the sequence's context, ultimately improving its prediction accuracy.

However, due to multiplicative effects, RNNs face challenges with exploding/vanishing gradients. Gers et al. (2000) report RNNs are generally limited to looking back in time for approximately 5-10 timesteps. Because of this limitation, traditional *Vanilla RNNs* are rarely used. Nevertheless, more sophisticated architectures derived from RNNs, like Long Short-Term Memory networks, effectively address these challenges and remain relevant.

3.3.3 Long Short-Term Memory Networks

Long Short-Term Memory networks (LSTMs), introduced by Hochreiter and Schmidhuber (1997), are an advanced type of RNNs designed to address the challenges of exploding and vanishing gradients in standard RNNs. Unlike traditional RNNs, which struggle with learning long-term dependencies due to gradient issues, LSTMs are adept at retaining information over longer sequences. Staudemeyer and Morris (2019) reports up to 1,000 steps, making them suitable for complex tasks such as time series forecasting.

The constant error carousel is a fundamental concept in LSTM networks, addressing the vanishing gradient problem. The constant error carousel is achieved through its unique cell state structure. The cell state in an LSTM is a longitudinal component that carries relevant information throughout the sequence of the network, enabling the retention and selective modification of information over long periods. It carries the network's memory, allowing information to be passed along unchanged if necessary (Staudemeyer & Morris, 2019). This is done by the LSTM's gated structure, which regulates the information added to or removed from the cell state.

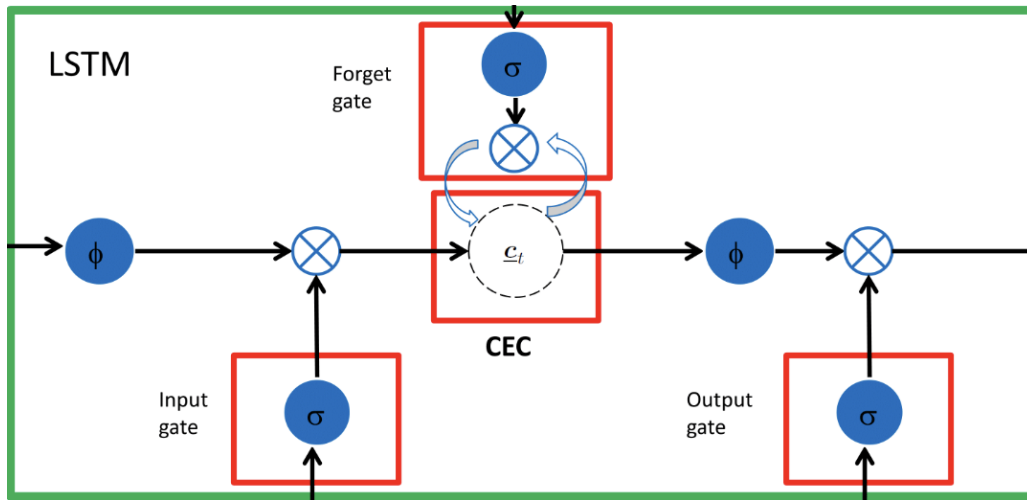


Figure 2: Illustration LSTM Memory Block (Su & Kuo, 2018)

At the core of LSTM's architecture are memory blocks, depicted in Figure 2, which replace the standard neurons found in RNNs. These memory blocks are equipped with the following arrangement of gates: (1) the forget gate, (2) the input gate, and (3) the output gate. The forget gate is responsible for deciding which information to discard from the block's memory, effectively "forgetting" irrelevant data. The input gate controls the addition of new information to the memory, updating the cell state with inputs. The output gate then determines the next hidden state, considering the current input and the block's memory (Staudemeyer & Morris, 2019).

LSTMs use a mixed learning method that combines different training approaches for better results. Importantly, forget gates decide what information to keep or remove, focusing on relevant data. In the backward step, the network uses BPTT to adjust weights based on learning from errors (Staudemeyer & Morris, 2019). Although LSTMs are good at learning from long data sequences and avoid issues with vanishing and exploding gradients, they require a lot of computing power and can be complex to fine-tune.

3.3.4 Convolutional Neural Networks

AlexNet by Krizhevsky et al. (2012) marked a breakthrough in computer vision. This neural network was made possible with convolutions. The origins of *Convolutional Neural Networks* (CNN) go back to 1989 when it was first introduced by Yann LeCun and his colleagues in *Handwritten Digit Recognition with a Back-Propagation Network*. Despite the innovation of the transformer, which now can outperform CNNs in computer vision tasks, CNNs remain one of the most influential predecessors. They are still used in various applications beyond computer vision (Deininger et al., 2022).

At the heart of CNNs are convolutional layers. A *convolution* is an operation used to combine an input with a filter, also known as a *kernel*, to produce a feature map. This operation highlights patterns or features in the input data. In a discrete convolution, the kernel slides over the input data, performing element-wise multiplications with the covered part of the input at each position. The results of these multiplications are summed to produce a single output value, forming the output feature map (Raschka et al., 2022).

For example, consider a two-dimensional input matrix X , representing time series data, and a kernel matrix W . The convolution operation can be expressed as:

$$Y = X * W \rightarrow Y_{i,j} = \sum_{k1=-\infty}^{+\infty} \sum_{k2=-\infty}^{+\infty} X_{i-k1, j-k2} W_{k1, k2} \quad (46)$$

Here, X is the input matrix, and W is the kernel matrix. Infinite bounds are not used in practical applications; padding is applied. Padding involves adding a finite number of zeros around the input matrix, ensuring that the indices remain within a valid range (Raschka et al., 2022). This allows the convolution to handle edge values appropriately and maintain the spatial dimensions of the input data.

In addition to convolutional layers, CNNs often incorporate pooling layers, which help to reduce the dimensionality of the feature maps. Pooling layers perform a down-sampling operation along the spatial dimensions of the input, which simplifies the information and reduces the computational load

for subsequent layers (Raschka et al., 2022). This operation helps detect certain features invariant to scale as it introduces local invariance and increases robustness to noise (Raschka et al., 2022).

This operation can be particularly useful for time series data. By applying a convolution across a time series, the model can detect patterns or features, like trends or seasonal effects, over sliding windows of the series, making it valuable for forecasting tasks.

3.4 Transformer-Based Models

In deep learning, CNNs and LSTMs have historically dominated the scene, particularly for processing sequence data. However, a shift occurred when Vaswani et al. (2017) from Google Brain and Google Research released *Attention Is All You Need*, which introduced the *transformer model*. The model has become a cornerstone in deep learning, particularly in NLP, gaining widespread adoption among tech giants such as Google, OpenAI, Microsoft, and Nvidia. Transformer-based models encompass a wide range of architecture that utilize the transformer mechanism for processing data. One notable subset of these models is OpenAI's *Generative Pre-trained Transformers* (GPTs).

Before transformers, RNNs, including LSTM units and Gated Recurrent Units (GRUs), were the preferred solutions for sequential data. These models, however, faced limitations in their ability to process information from a long series of data due to their fixed reference window, which restricted how far back they could look to make predictions. While LSTMs mitigated some of these challenges, they did not completely overcome them. Additionally, the computational demands for RNN-based models scaled exponentially with task complexity, posing practical limitations for their application in real-world scenarios.

Transformers introduced a revolutionary approach with their theoretically infinite reference window, allowing them to consider all available data points when making predictions. This capability marked significant progress in NLP and showed promise for application in time series forecasting despite the differences between these fields. Moreover, the transformer architecture facilitated parallel processing during training, drastically reducing the computational resources required. This innovation addressed the scalability issues previous neural network models faced and opened new gateways for research and application in various domains beyond NLP.

3.4.1 Architecture

The Transformer model's architecture comprises an *encoder* and a *decoder*. Typically, the input of a neural network is called a *token*. While practitioners and scholars seem to lack agreement on this

for time series, we will use this term. The encoder processes the input token, encoding it into a high-dimensional space where relationships are captured through layers of self-attention and feed-forward networks. The decoder then uses this encoded information to generate future values in the series, leveraging its self-attention layers to maintain internal coherence and encoder-decoder attention layers to focus on relevant input parts. The training of a transformer model is much the same as that of other neural networks, where gradient descent plays a vital role.

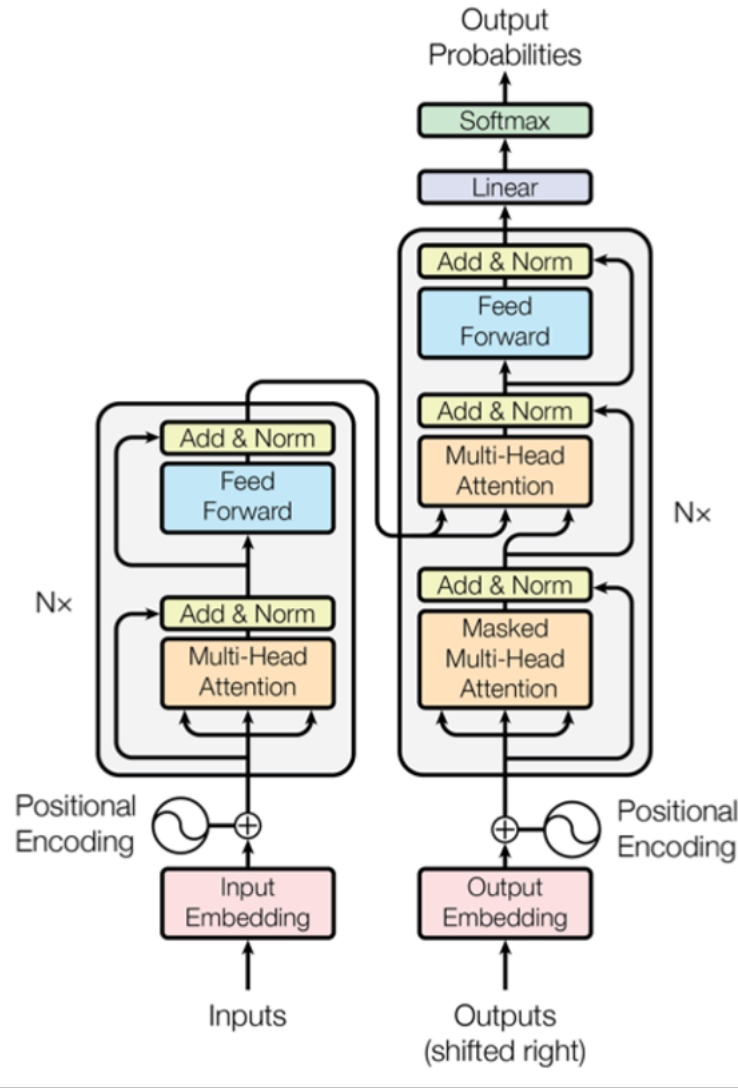


Figure 3: Illustration of the original Transformer model from Vaswani et al. (2017).

Positional Encoding and Embedding Because transformers do not process tokens sequentially, such as an RNN or CNN, the model needs to understand the order in a different manner. To solve this, Vaswani et al. (2017) uses *embedding* and *positional encoding*.

Embedding is a pre-processing step that translates raw data into a numerical representation suitable for neural network models. In NLP, embedding converts textual elements, such as words or phrases, into vector representations. Techniques like Word2Vec are widely employed to achieve this trans-

formation. These vectors capture the semantic relationships between words, allowing the model to determine the meaning of a word based on its context. For time series data, this would extend to time series attributes such as seasonality, cycles, and stationarity.

Positional encoding supplements the embedding vectors with additional information that captures the relative position of each element within the sequence. This additional information allows the model to retain the sequential information and understand how the elements relate. In Vaswani et al. (2017), they use sine and cosine functions for positional encoding:

$$PE_{(pos,2i)} = \sin\left(\frac{pos}{10000^{2i/d_{model}}}\right) \quad (47)$$

$$PE_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{2i/d_{model}}}\right) \quad (48)$$

Where pos is the position, i is the dimension, and d_{model} represents a dimension of the models' embeddings, meaning that if a token is represented by a 512-dimensional vector, d_{model} would equal 512. Note the difference in the subscript as the *sine* function is applied to those at even positions, and *cosine* is applied to those at odd positions.

Encoder and Decoder The encoder is the first component of a transformer model and is responsible for processing and understanding the input sequence. It consists of multiple encoding layers containing a *self-attention mechanism* (discussed in detail below) and a feed-forward network. The feed-forward network is a simple neural network that is used to refine the representation of the input token further, enabling the model to capture non-linear relationships and transformations in the token.

In the transformer architecture, depicted in Figure 3, each of the sub-layers within a transformer block, which includes the multi-head attention layer and the position-wise feed-forward network, is followed by *layer normalization*, a technique introduced by Ba et al. (2016). Specifically, the output of each sub-layer passes through a residual connection that adds the sub-layer input to its output, and the layer normalization then normalizes this combined output. The layer normalization ensures that each component output is standardized before passing to the next layer or sub-layer, contributing to the model's effectiveness and stability. Vaswani et al. (2017) highlights that normalization helps stabilize the learning process by ensuring consistent scale and mean of activations across the network's depth.

Attention in Transformers Central to transformer models is the mechanism of attention. Bahdanau et al. (2016) first introduced in their work to improve the performance of machine translation. The mechanism involves computing alignment scores between different points in the data. These scores indicate the relevance or importance of each data point when predicting a future value. The model then creates a context vector for each predicted point, a weighted sum of the data, with weights derived from the alignment scores. The alignment scores are typically calculated using the hidden states from a recurrent neural network or another model framework, reflecting the relationship between each point.

Contrastive to models such as RNN, which relies on sequential processing, Vaswani et al. (2017) reports transformers can focus on different parts of the input sequence in parallel, irrespective of their position, using self-attention. When applied to time series, it enables each point in the series to attend to all other points, in contrast to traditional time series models that often rely on a fixed window of recent data. Processing all points in a time series in parallel offers a significant efficiency advantage.

The model generates three vectors for each point: a *query* vector, a *key* vector, and a *value* vector. These are produced from the same data point but are transformed differently. One can consider these in the following way: *queries* is an inquiry to find some specified information stored in *keys*, and when matched, they return a *value*. While this is highly simplified, it encapsulates the logic of the attention mechanism, which for time series is finding points relevant to the value at hand. To do this, Vaswani et al. (2017), uses *Scaled Dot-Product* (multiplicative) *Attention*. This is computed using the following formula:

$$Attention(Q, K, V) = \text{SoftMax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \quad (49)$$

In practice, the *queries*, *keys*, and *values* are matrices. QK^T represents the matrix product of the *query* matrix and the transpose of the *key* matrix. The scaling, dividing by $\sqrt{d_k}$ (the square root of the vector length), is applied to avoid issues with unsuitably small gradients from the SoftMax. An optional feature of the attention mechanism is *masking*. Vaswani et al. (2017) uses masking in this context to prevent the model from accessing future values of the input sequence when computing the attention score for i , ensuring the prediction only has “known” values. This step is done before the SoftMax function is applied. The output is a sum weighted by the *value* matrix, describing the relationship and dependency of the order, reflecting how the different points pay “attention” to each other. In practice, *Multi-Head Attention* is used, where this process is done multiple times in parallel to capture information from different positions. This translates to:

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) W^O \quad (50)$$

Where $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$ and W_i^Q, W_i^K , and W_i^V are weight matrices and W^O is the weights of the output. The number of heads, or attention layers, is a specified parameter. Because the outputs of all attention heads are concatenated and then multiplied with the output weights, the model can draw insights from all heads.

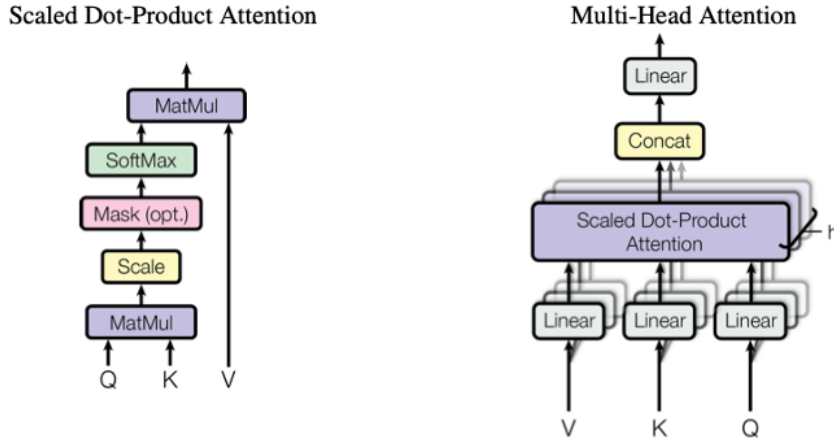


Figure 4: Illustration of Multi-Head Attention (Vaswani et al., 2017)

3.4.2 Inference

After the transformer is trained, it is ready for inference. The entire input sequence is provided to the encoder, which processes it all simultaneously and generates a series of encoder outputs. These outputs capture the contextual relationships within the input sequence. Umar Jamil (2023) explains that the inference process begins in the decoder with an initial start-of-sequence (SOS) token. This token and the encoder outputs are fed into the decoder. The decoder then generates the first output token by processing the combination of the SOS token and the encoder outputs through its layers.

This output from the decoder is transformed by a linear layer followed by a SoftMax function, which translates it into a probability distribution over the possible following tokens. The token with the highest probability is selected as the output. The decoder takes the previously generated tokens, and the encoder outputs them as input for each subsequent token. This process is repeated, with the decoder generating one token at a time until it produces an end-of-sequence (EOS) token or completes the sequence. Each decoder output is subject to the same linear and SoftMax transformation to determine

the next token in the sequence. This sequential generation continues until the entire output sequence is produced.

3.4.3 Transfer Learning

In traditional machine learning, the training and test data are expected to be very similar, with the same input feature space and data distribution. However, a limited supply of training data can occur due to the data being limited, rare or expensive to collect and label. According to Weiss et al. (2016), pp. 1–2, this introduced the need for *transfer learning* as a capacity to apply knowledge from a related source domain to the target domain. This is particularly important for foundation models, like our transformer-based models, as they rely on generalizing across domains, even on new datasets on which they have not been trained. As illustrated in Figure 5, traditional machine learning has a one-to-one relationship between each task and learning system, while transfer learning aims to utilize knowledge from multiple source tasks in a learning system on a new target task.

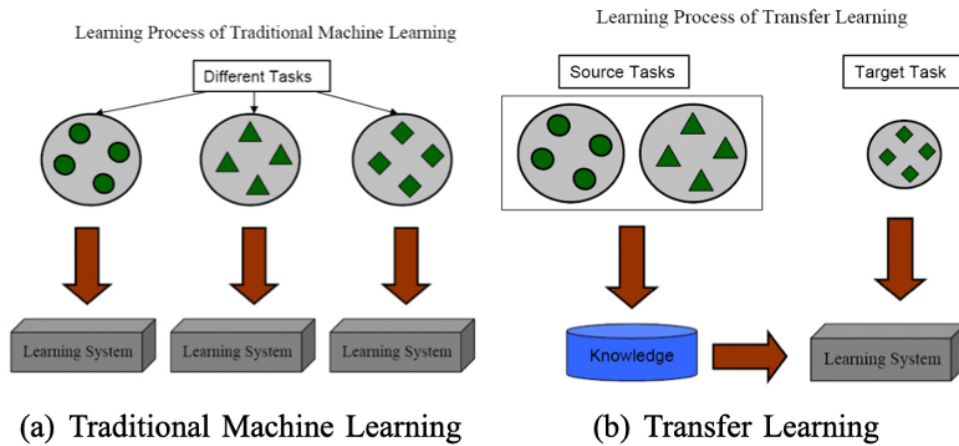


Figure 5: Illustration of Transfer Learning (Pan & Yang, 2010)

Foundation models aim to use knowledge from pre-trained data to handle various tasks. By training on a vast amount of diverse data, these models learn patterns and relationships that can be applied to different areas. This is helpful when labeled data is hard to get or expensive because it allows the models to work well on new tasks with minimal extra training. Transformers can be trained with supervised and unsupervised learning. In supervised learning, the model is trained on labeled data, while in unsupervised learning, the model is trained on unlabeled data.

However, a limitation of foundation models is the training. Foundation models require extensive pre-training datasets and are computationally intensive, which can make training time-consuming, often requiring days or weeks, despite efforts to optimize the training process (Ivanov et al., 2021; Yang

et al., [2022](#)).

3.4.4 Transformers for Time Series Forecasting

Recent research in the field of transformer models for time series forecasting highlights efforts to modify the architecture to suit this task better. This research has led to several suggested architectures. Consider how a transformer model processes text to motivate these changes in the architecture. In textual data, the model needs to grasp the relationships between words to understand longer contexts, and it achieves this through the attention mechanism. For example, consider the sentence: *"I wish your help in the project I am working on"*. It still conveys the same meaning if we rephrase it to: *"I am working on a project, and I need your help"*. Ideally, a transformer model should understand both sentences equally well despite the different word order. However, with time series data, the order of the data points is crucial. Imagine a fictitious series representing daily temperatures over a week: $[21, 23, 20, 22, 24, 25, 21]$. Reordering this to $[25, 21, 23, 20, 24, 22, 21]$ would distort the entire trend and lead to incorrect insights.

In other words, while transformers can handle word reordering in text data, reordering can destroy meaning in time series data. Adapting the transformer architecture aims to capture temporal dependencies more precisely. This objective has led to innovative approaches within the transformer framework, such as the Autoformer by Wu et al. ([2021](#)).

Autoformer advances the architecture by incorporating an automatic decomposition mechanism that separates the series into trend and seasonal components before processing. Wu et al. ([2021](#)) report that this decomposition allows the model to focus on the underlying trend by applying a series-wise auto-correlation window, improving the forecasting accuracy for time series with strong seasonal patterns. Additionally, Autoformer redesigns the self-attention framework to integrate a frequency-enhanced attention mechanism that aids in distinguishing between different frequencies of time series data, thereby enhancing the capture of temporal dependencies.

Another proposed architecture to better suit time series forecasting is the Informer model by Zhou et al. ([2020](#)). This model addresses the inefficiencies in standard self-attention mechanisms, which scale poorly with longer time sequences due to quadratic computational complexity. Zhou et al. ([2020](#)) introduces a sparsity mechanism in the attention module, termed "ProbSparse Self-Attention". This attention mechanism selectively activates the most informative parts of the data, reportedly reducing the computational load without sacrificing performance. This makes it suitable for processing very long time series efficiently. Furthermore, the Informer employs a distilling operation that further

compresses the sequence length during the attention process, enhancing its ability to capture temporal dependencies.

These adaptations aim to address the main limitation of the original transformer design: its inability to maintain order in the data, which is crucial for accurate time series forecasting. Refining and adjusting the architecture makes the transformer model aims to make it more adept at time series forecasting.

On the other hand, Zeng et al. (2022) critically examines the applicability of transformers to time series forecasting. While natural language tasks rely on the sequential order of words, time series forecasting demands precise handling of time series—a challenge for transformers due to their self-attention mechanism, which is relatively insensitive to the order of inputs, potentially leading to the loss of time series attributes. Zeng et al. (2022) demonstrate that despite their sophisticated architecture, transformers can underperform compared to simpler models, such as a basic one-layer linear model, particularly in long-term time series forecasting scenarios. This juxtaposition raises questions about the practical effectiveness of transformers in time series contexts, calling for further adaptation in the field.

Moreover, Zeng et al. (2022) explore various aspects of transformer design that impact their ability to extract temporal dependencies. Their studies show that modifications to the self-attention mechanism intended to capture time series attributes better do not necessarily overcome its fundamental limitations. They conclude that the effectiveness of transformers for long-term time series forecasting is overstated and recommend reassessing their use in time series analysis, especially in cases where accurate modeling and understanding of time-related patterns are crucial.

These findings underscore the need for continued development and evaluation of time series forecasting models. They question the reliance on complex architectures like transformers when simpler, more interpretable models may be sufficient.

3.5 TimeGPT

TimeGPT is a foundation model for time series data developed by the company Nixtla and was introduced by Garza and Mergenthaler-Canseco (2023). Unlike open-source projects, TimeGPT is commercially driven, which results in limited technical documentation. Nevertheless, this section aims to provide an overview of the model’s architecture, training, how it creates forecasts and other capabilities. The version of TimeGPT relevant to this thesis is TimeGPT-1.

3.5.1 Model Architecture

TimeGPT's architecture, depicted in Figure 6, is similar to that of Vaswani et al. (2017). However, some key elements differ. Naturally, it does not output probabilities from a SoftMax function; instead, the last layer is linear to output predicted continuous values. Furthermore, where the original transformer uses feed-forward networks, TimeGPT has replaced this with convolutional neural networks. However, it is not specified in any way how this is used or why. Garza and Mergenthaler-Canseco (2023) also underlines that it is not based on any existing LLM and that it is specialized to handle time series data with the goal of minimizing forecasting error. Apart from this, Garza and Mergenthaler-Canseco (2023) does not reveal any other significant differences in the model architecture.

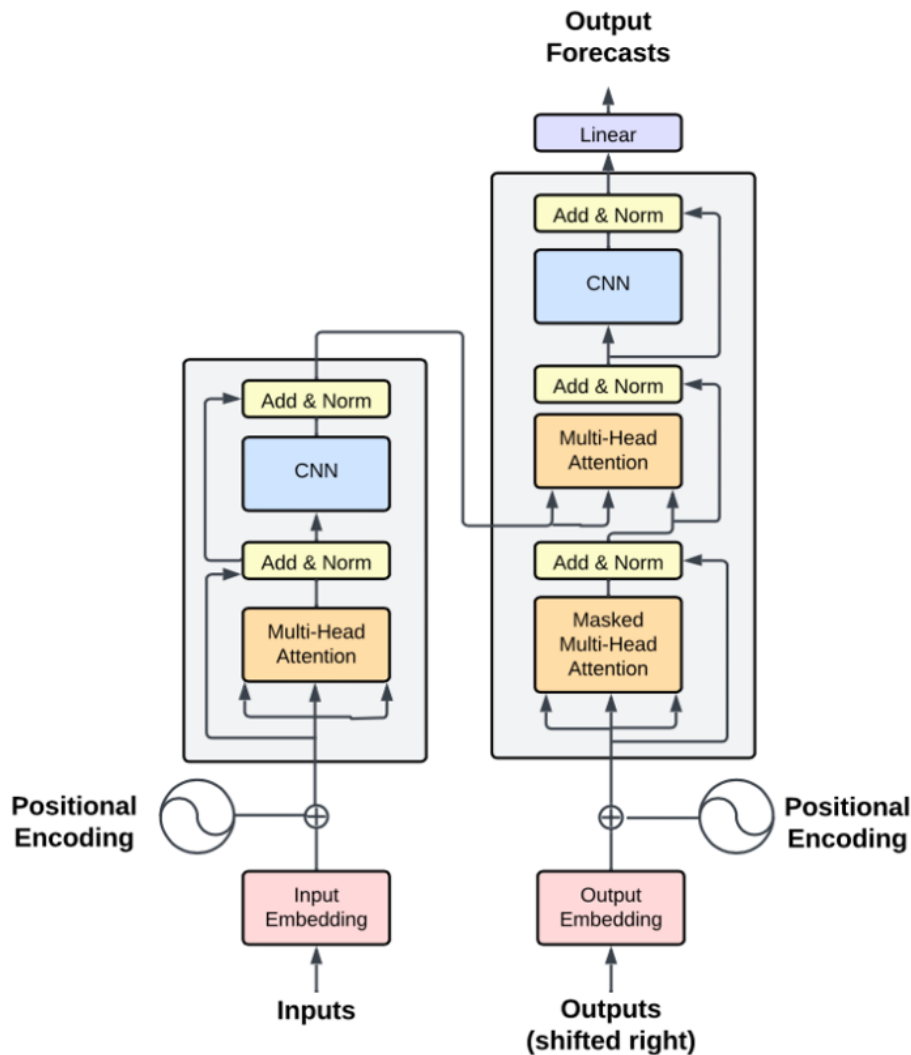


Figure 6: Model architecture of TimeGPT (Garza & Mergenthaler-Canseco, 2023)

3.5.2 Training

Garza and Mergenthaler-Canseco (2023) claims that TimeGPT was trained on the largest publicly available time series dataset as of 5. November 2023 with more than 100 billion observations. The data is reportedly from various domains such as finance, health, industry, sales, IoT sensors, energy, transport, weather, etc. It is mentioned that the data contains several different forms of seasonality and cycles, representing non-stationary real-world data. Regarding preprocessing, it is only noted that it was formatted to a standard and that missing values were filled in. It is unclear how this was done, whether by imputing the mean, median, or another method. In addition, no statistics such as time frequency about the data is reported.

In the training of the model, they mention that hyperparameters such as the learning rate and batch sizes were explored and tampered with extensively. In addition, they do mention that it was trained using the Adam optimizer with a decaying learning rate (described in 3.3.1).

3.5.3 Approach to Time Series Forecasting

TimeGPT's approach to forecasting involves reading time series data as a sequence of tokens. With a target variable and the possibility of including events and additional variables, TimeGPT can process input data to create forecasts (Nixtla, n.d.).

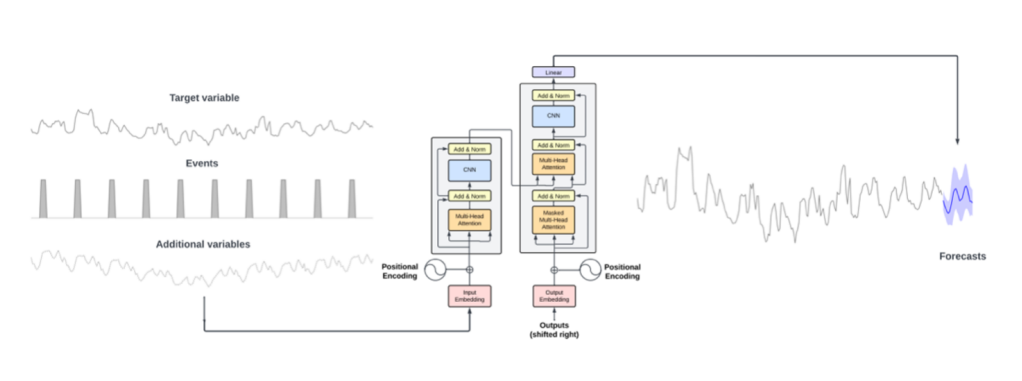


Figure 7: The Forecasting Pipeline of TimeGPT (Garza & Mergenthaler-Canseco, 2023)

As the model ingests time series data, it processes the information in chunks, transforming raw numerical inputs into discrete tokens that serve as input. As mentioned earlier, TimeGPT utilizes a structure of CNNs and multi-head attention mechanisms, supposedly allowing it to focus on different parts of the time series and appropriately weigh their significance when making predictions.

The model utilizes transfer learning with its pre-trained foundation. As explained earlier, the model is exposed to diverse time series datasets across various domains. This exposure enables the model

to transfer the knowledge acquired during training to create forecasts for new, unseen time series. The pre-trained foundation aims to equip TimeGPT with an understanding of time series behaviors to empower the model to adapt to new time series data quickly. Garza and Mergenthaler-Canseco (2023) claims it allows TimeGPT to generalize across different fields without extensive retraining or fine-tuning, although fine-tuning can enhance model performance.

3.5.4 Available Methods and Capabilities

TimeGPT offers several capabilities, such as anomaly detection and the inclusion of exogenous variables. In this thesis, we are primarily interested in its zero-shot forecasting ability. Additionally, we conducted a small experiment using fine-tuning. For the fine-tuning experiment, please see [Appendix D](#).

3.6 Chronos

Ansari et al. (2024) introduced Chronos, a set of multiple pre-trained time series forecasting models. It is inspired by the success of large language models (LLMs) and uses existing transformer-based language model architecture for time series. Chronos is open source, as it is freely available and collaborative to the public. Amazon Web Services developed Chronos in collaboration with UC San Diego, the University of Freiburg, and Amazon Supply Chain Optimization Technologies.

3.6.1 Model Architecture

Chronos follows a minimalistic approach, utilizing an existing model architecture, the Text-to-Text Transfer Transformer (T5) by Raffel et al. (2020). The only difference between Chronos and the original T5 is the vocabulary size, which is the number of tokens a model recognizes and utilizes to process and generate, resulting in fewer parameters. T5 has no time-series-specific features or design. The authors behind Chronos questioned whether any design modifications were essential for effective forecasting. Thus, Ansari et al. (2024) developed Chronos using a language modeling framework with minimal adaptations for time series forecasting.

The architecture for T5 is close to that of Vaswani et al. (2017). The main differences in architecture are that T5 apply a simplified layer normalization, placing this layer before the feed-forward and attention components and utilizing a simpler scheme for positional encoding. The impact of the modifications made from the original transformer architecture has not yet been fully explored. However, given that these changes were considered minor in the research conducted by Raffel et al. (2020), it is reasonable to assume that they will also be insignificant for this thesis.

3.6.2 Training

The Chronos models are trained on the base of a large set of publicly available datasets spanning multiple domains. This totals 28 datasets consisting of 893,208 time series. There is also variety in the time intervals for the time series, ranging from five-minute intervals to monthly intervals. The authors used data augmentation techniques to artificially enhance the training data to create a more robust foundation for the model. For the Chronos models, data augmentation includes (1) Time Series Mixup and (2) Synthetic Data Generation. Time Series Mixup (TSMix) randomly samples time series from the training data and scales them before combining them into a new time series. This improves pattern diversity. However, as the authors do not deem it sufficient to create a generalist time series model, they further generate synthetic data using Gaussian processes with the method KernelSynth. This method uses the underlying structures of the training data to create new time series'. This resulted in the Chronos models being trained on 10M TSMix augmentations and 1M synthetic time series generated with KernelSynth. Both TSMix and KernelSynth are innovations introduced by Ansari et al. (2024).

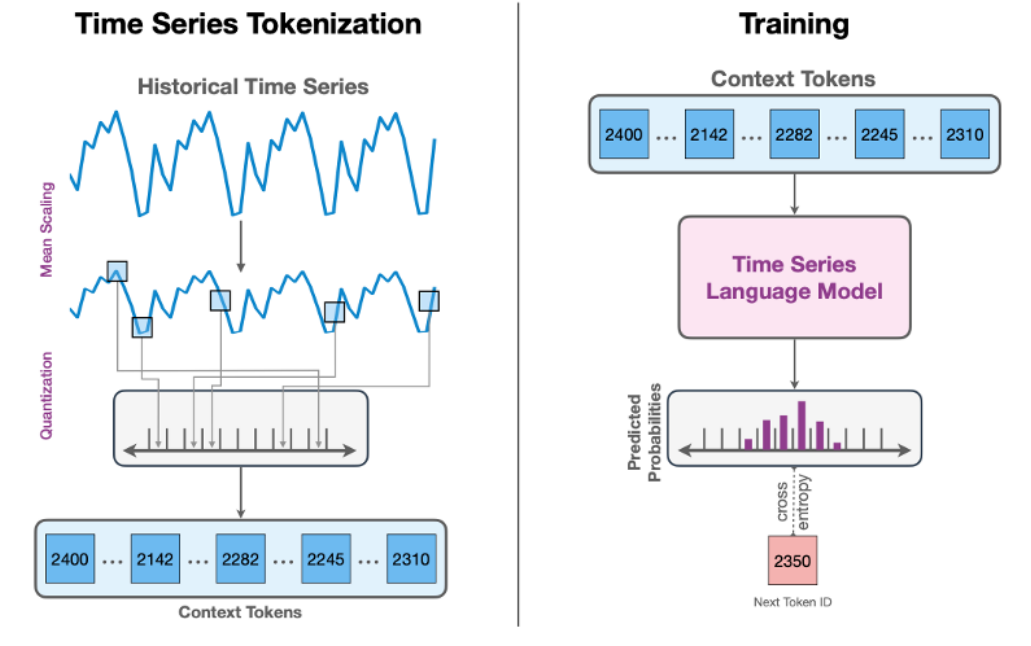


Figure 8: High-level Depiction of the Training of Chronos (Ansari et al., 2024)

Data from the M4 Competition is used to pretrain the Chronos models and conduct an in-domain evaluation. However, data from the M3 Competition is only used for zero-shot evaluation by Ansari et al. (2024), meaning that the model is not pre-trained on the data used in this thesis.

3.6.3 Approach to Time Series Forecasting

As previously mentioned, the T5 architecture was designed to handle text. It processes text by breaking it down into tokens and generating responses with tokens that are contextually relevant to the input. Through its training, it has learned to identify appropriate tokens. Consequently, the model learns tokens rather than the actual text data. The rationale behind Chronos is that when the tokens are adapted to handle time series data, the model can still produce appropriate tokens that yield meaningful outputs for time series forecasts. Ansari et al. (2024) reports that Chronos uses a categorical distribution to model the observations and performs regression through a classification approach known as regression-via-classification.

Ansari et al. (2024) describes that historical time series data is scaled and then quantized into tokens, which serve as the model's input. The scaling normalizes the data, and quantization converts numerical values into discrete tokens, which are more suitable for language models. For generating forecasts, the trained model samples tokens autoregressively. That means it uses its previous output as the input for each step to predict the next token. These tokens are then mapped back to numerical values, which are the actual forecasts. By sampling multiple trajectories, the model builds a probabilistic forecast, capturing the uncertainty in the predictions.

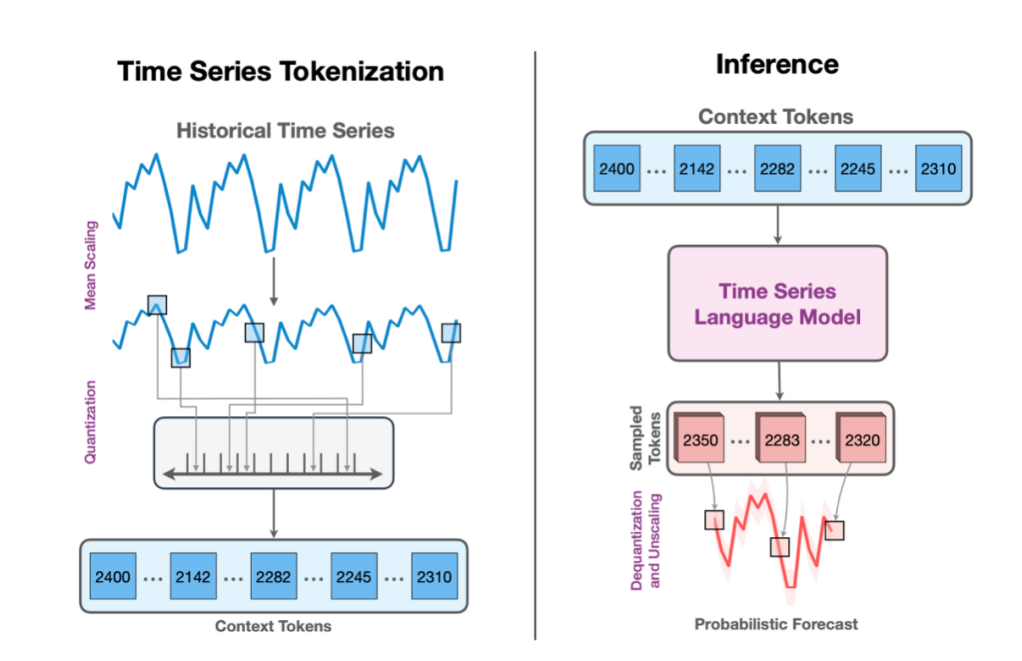


Figure 9: High-level Depiction of Inference of Chronos (Ansari et al., 2024)

3.6.4 Available Methods and Capabilities

The Chronos models are the five models in Table 1. The only difference between them is the parameter sizes. Ansari et al. (2024) demonstrates that the bigger the model is, the better the result in their analysis. However, the largest model is significantly more computationally expensive. This is also why Chronos does not include larger models, as they would become impractical in real-world applications.

Table 1: The Five Models of Chronos (Ansari et al., 2024)

Model	Parameters	Based On
chronos-t5-tiny	8M	t5-efficient-tiny
chronos-t5-mini	20M	t5-efficient-mini
chronos-t5-small	46M	t5-efficient-small
chronos-t5-base	200M	t5-efficient-base
chronos-t5-large	710M	t5-efficient-large

3.7 MOIRAI

Woo et al. (2024) introduced the *Masker Encoder-based Universal Time Series Forecasting Transformer* (MOIRAI), a foundation model for time series forecasting, drawing inspiration from LLMs and vision transformers. Moirai utilizes a unified transformer-based architecture tailored for time series data. It is an open-source initiative developed collaboratively by Salesforce AI Research and the School of Computing and Information Systems at Singapore Management University.

3.7.1 Model Architecture

Moirai aims to handle time series data effectively by using a Masked Encoder architecture. This setup includes several input and output layers that can handle the different frequencies and types of data found in time series. These layers aim to help the model learn effectively from data that varies in speed and complexity.

Unlike typical transformers that handle single-dimension data, Moirai introduces *Any-variate Attention* that treats all variates of a multivariate time series as a single sequence (Woo et al., 2024). This supposedly allows the model to process the full context of multivariate interactions. In the architecture, the basic transformer structure is maintained, but the data handling is adapted in an effort to meet the needs of universal time series forecasting.



Figure 10: High-level architecture of Moirai (Woo et al., 2024)

3.7.2 Training

Moirai is trained on the Large-scale Open Time Series Archive (LOTSA), containing over 27 billion observations across various domains (Woo et al., 2024).

Table 2: Key Statistics Frequency and Domain of the LOTSA Data (Woo et al., 2024)

	Energy	Transport	Climate	CloudOps	Web	Sales	Nature	Econ/Fin	Healthcare
# Datasets	30	23	6	3	3	6	5	23	6
# Obs.	16,358,600,896	4,900,453,419	4,188,011,890	1,518,268,292	428,082,373	197,834,339	28,547,647	24,919,596	1,594,281
%	59.17%	17.73%	15.15%	5.49%	1.55%	0.72%	0.09%	0.10%	0.01%

	Yearly	Quarterly	Monthly	Weekly	Daily	(Multi) Hourly	(Multi) Minute-level	(Multi) Second-level
# Datasets	4	5	10	7	21	31	25	2
# Obs.	873,297	2,312,027	11,040,648	18,481,871	709,017,118	19,875,993,973	7,013,949,430	14,794,369
%	0.003%	0.008%	0.040%	0.067%	2.565%	71.893%	25.370%	0.054%

This diverse dataset helps Moirai learn a wide range of time series patterns, from economic indicators to climate data. Instead of just predicting single values, the model learns to predict the entire range of possible future values. This is done by training it to optimize a mix of parametric distributions, which prepares the model for both exact and probabilistic forecasts. The small model is trained for 100,000 steps, while the base and large models are trained for 1,000,000 steps. The training uses the Adam algorithm with weight decay (described in 3.3.1).

Moirai uses techniques like random sampling of context lengths and prediction horizons during training. This approach adds flexibility and robustness to the model’s forecasting abilities.

However, there is a limitation to note for our thesis: Moirai has been trained on LOTSA, which includes data from the M3 Competition. This overlap will be considered since our thesis evaluates foundation models for time series using the M3 dataset.

3.7.3 Approach to Time Series Forecasting

Moirai does not generate point estimates when producing forecasts but predicts a distribution of future values. This is achieved by optimizing a mixture of parametric distributions, which allows the model to account for the inherent uncertainty in time series data Woo et al. (2024). By focusing on the distribution of possible future outcomes, Moirai provides forecasts that include both deterministic and probabilistic elements, reportedly making it versatile for various forecasting requirements.

3.7.4 Available Methods and Capabilities

The model comes in three versions to meet different computation and detail needs: Moirai-Small, Moirai-Base, and Moirai-Large. These versions have 14 million, 91 million, and 311 million parameters, respectively. Each version is designed to balance performance and computational efficiency, with the larger models offering more complexity and potentially higher accuracy in forecasting tasks across diverse datasets.

3.8 Model Evaluation Metrics

It is important to choose an appropriate metric when evaluating a model. Choosing these metrics is, unfortunately, not straightforward as there are many to choose from, and the different metrics serve different purposes. This subsection covers the ideas behind the evaluation metrics used in this thesis.

The metrics aim to measure the forecast accuracy by summarizing, in different ways, the *forecast errors*, where the forecast errors are the difference between observed values and the forecast. Errors do not refer to a mistake but rather the unpredictable part of an observation and can be written as e_{T+h} for the training data $\{y_1, \dots, y_T\}$ and the test data $\{y_{(T+1)}, y_{(T+2)}, \dots\}$:

$$e_{T+h} = y_{T+h} - \hat{y}_{T+h|T} \quad (51)$$

It is important to note that the forecast errors are not the same as residuals. What differentiates forecast errors and residuals is: (1) residuals are calculated based on the observed values and the forecast on the training data, while forecast errors are calculated on test data, and (2) residuals are based on forecasting one time period ahead of each step while forecast errors can be based on the prediction of several future time periods (Hyndman & Athanasopoulos, 2021).

Forecast errors can be further extended to find the *percentage* error. This has the advantage of being unit-free and is therefore useful to compare forecast performance for different data sets (Hyndman &

Athanasopoulos, 2021). The percentage error p_t is given as:

$$p_t = \frac{e_t}{y_t} * 100 \quad (52)$$

This is to convert the forecast error e_t , the unit difference, to a percentage. Earlier, we defined e_{T+h} to represent the forecast for h periods ahead of time T . However, for simplicity and because the focus will shift solely to what would be the test data, we will use just time t in subsequent discussions with n samples of forecast errors.

An alternative way of measuring is to base the calculation on *relative errors*. This lets us compare the forecast error of a method with a benchmark. The relative error r_t is:

$$r_t = \frac{e_t}{e_t^{\text{benchmark}}} \quad (53)$$

In this formula, we find the relative error r_t by dividing the forecast error e_t by the forecast error of a benchmark method $e_t^{\text{benchmark}}$. This benchmark method is usually a *random walk*, meaning the forecast is equal to the last observation (Hyndman & Koehler, 2006, p. 683).

3.8.1 sMAPE

The *Symmetric Mean Absolute Percentage Error* (sMAPE) is a central metric to the M3 Competition. To understand sMAPE, starting with the related metric Mean Absolute Percentage Error (MAPE) is useful as it originates from this metric from which sMAPE originates. With the definition of p_t from (52), MAPE is given as:

$$\text{MAPE} = \frac{1}{n} \sum_{t=1}^n |p_t| \quad (54)$$

As shown, MAPE calculates the mean of the percentage error p_t in absolute values. The MAPE has the disadvantage of becoming infinite or undefined when y_t is zero and extremely skewed when y_t is close to zero. Another disadvantage is that the MAPE puts a heavier penalty on positive errors than on negative errors, which led to the use of sMAPE, with a change in the denominator for MAPE to introduce “symmetry” (Hyndman & Koehler, 2006, p. 683). The formula used in the competition with forecast error e_t , observed values y_t and forecast \hat{y}_t is:

$$\text{sMAPE} = \frac{1}{n} \sum_{t=1}^n \frac{|e_t| * 100}{(y_t + \hat{y}_t)/2} = \frac{1}{n} \sum_{t=1}^n \frac{|e_t| * 200}{(y_t + \hat{y}_t)} \quad (55)$$

The denominator, $|y_t| + |\hat{y}_t|$, aims to add symmetry by balancing the error in relation to the size of both the observed and forecasted values. This means that the forecast error in an absolute value will be divided by the average of the observed and forecasted values to create a more stable error metric.

With sMAPE, problems with small values for y_t are less severe as the forecast value \hat{y}_t is added to the denominator. However, when y_t is close to zero, \hat{y}_t is also likely to be close to zero if the forecast is accurate (Hyndman & Koehler, 2006, p. 683). This means the measure can still involve a division by a number close to zero, potentially making it unstable. Furthermore, sMAPE can take negative values, as the denominator is the average of the forecast and actual values, not their absolute values, making it not truly a metric of “absolute” percentage errors (Hyndman & Athanasopoulos, 2021). The denominator, however, is defined in absolute values for the M4-Competition (Makridakis et al., 2019). Additionally, the penalty for a low value of $\frac{|e_t|}{(y_t + \hat{y}_t)/2}$ is more costly compared to a higher value with the same y_t . This implies that a model consistently underestimating is penalized less severely than one that overestimates by the same absolute value, making the metric not as symmetric as its name might suggest (Goodwin & Lawton, 1999).

The sMAPE is also used to determine the **Average Ranking** in the competition, which is the average ranking of the sMAPE for each method for each forecast horizon (Makridakis & Hibon, 2000).

Although this metric is central in the M3 Competition and is widely used, Hyndman and Koehler (2006) argues that sMAPE should be avoided due to the mentioned issues. We will, therefore, not use this metric to the extent done in the M3 Competition to evaluate our findings.

3.8.2 MdAPE

Median Absolute Percentage Error (MdAPE) is one of the most used metrics based on percentage errors. It is given as:

$$\text{MdAPE} = \text{median}(|p_t|) \quad (56)$$

This rather simple metric returns the median percentage error p_t in absolute values. The MdAPE has the advantage, as mentioned earlier, of being metric based on percentage error. However, it has the same disadvantages as discussed for MAPE with being infinite or undefined for $y_t = 0$, being skewed

when close to zero and has asymmetry in penalties between positive and negative errors (Hyndman & Koehler, 2006, p. 683).

3.8.3 MdRAE

The *Median Relative Absolute Error* (MdRAE) is similar to MdAPE, but the percentage error p_t is exchanged for the relative error r_t :

$$\text{MdRAE} = \text{median}(|r_t|) \quad (57)$$

The MdRAE will, as the name reveals, be the median of the relative errors in absolute values. As mentioned, we find the relative error r_t by comparing two methods, often a random walk as the benchmark. In the M3-Competition, the seasonal naïve method is used as the benchmark, which is a variation of a random walk as covered in Section (3.2.1), Naïve Forecasting. However, this metric also has drawbacks. When the relative error to the benchmark $e_t^{\text{benchmark}}$ is close to zero, it introduces an instability, as with sMAPE, which is not ideal (Hyndman & Koehler, 2006, p. 683).

3.8.4 RMSE

The *Root Mean Square Error* (RMSE) is one of the most used scale-dependent accuracy metrics. It is scale dependent as it is only based on the forecast error e_t and is therefore incompatible for comparisons between series with different units (Hyndman & Athanasopoulos, 2021). The RMSE is calculated as:

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{t=1}^n e_t^2} \quad (58)$$

In the RMSE, all forecast errors e_t are squared, returning a positive number. The square root of the average of these numbers will be the value for this metric. As it sums the squared forecast errors, larger errors will have a bigger impact on the result than other metrics. It is also more sensitive to fluctuations in the forecast errors, in absolute values, especially if the errors are inconsistent (Willmott & Matsuura, 2005). This means that if there are big swings in how much the predictions are off, the RMSE will penalize this heavier than other metrics.

3.8.5 MASE

Hyndman and Koehler (2006) proposes to use scaled error in an accuracy metric *Mean Absolute Scaled Error* (MASE) as a better alternative than the previous metrics. As such, this metric will be central to our evaluation. The scaled error q_t is defined when benchmarking against seasonal naïve method as:

$$q_t = \frac{e_t}{\frac{1}{n-m} \sum_{t=m+1}^n |y_t - y_{t-m}|} \quad (59)$$

This is to make the metric independent of the scale of the data. The denominator represents the average error of a naïve one-step ahead prediction $y_t - y_{(t-1)}$ in absolute values. This means that the scaled error q_t measures forecast error normalized against the average of the forecast errors a naïve method would have. To then obtain the MASE, we simply average these scaled errors q_t in absolute values over the forecast period n :

$$\text{MASE} = \frac{1}{n} \sum_{t=1}^n |q_t| \quad (60)$$

While MASE might initially seem complex, it is a user-friendly indicator of forecast accuracy: a score below 1 indicates better performance than the naïve benchmark, around 1 means similar performance, and above 1 suggests poorer performance (Hyndman & Koehler, 2006, pp. 684–685).

4 Methodology

4.1 Introduction

In this section, we outline the methodology used to evaluate the performance of foundation models in zero-shot forecasting with the M3-Competition dataset. First, we describe the M3-Competition dataset, which includes a diverse collection of time series data across yearly, quarterly, and monthly frequencies. This dataset is widely used and provides robust benchmarks for comparison. Next, we detail the evaluation process. We use the Mean Absolute Scaled Error (MASE) score to measure forecast accuracy, comparing the foundation models to the Naïve2 benchmark. Finally, we explain how the foundation models—Chronos, TimeGPT, and Moirai—were used in this study. We describe their setup for zero-shot forecasting, the data preprocessing steps, and the systematic recording of results for analysis. Figure 11 shows a high-level depiction of the methodology of this thesis.

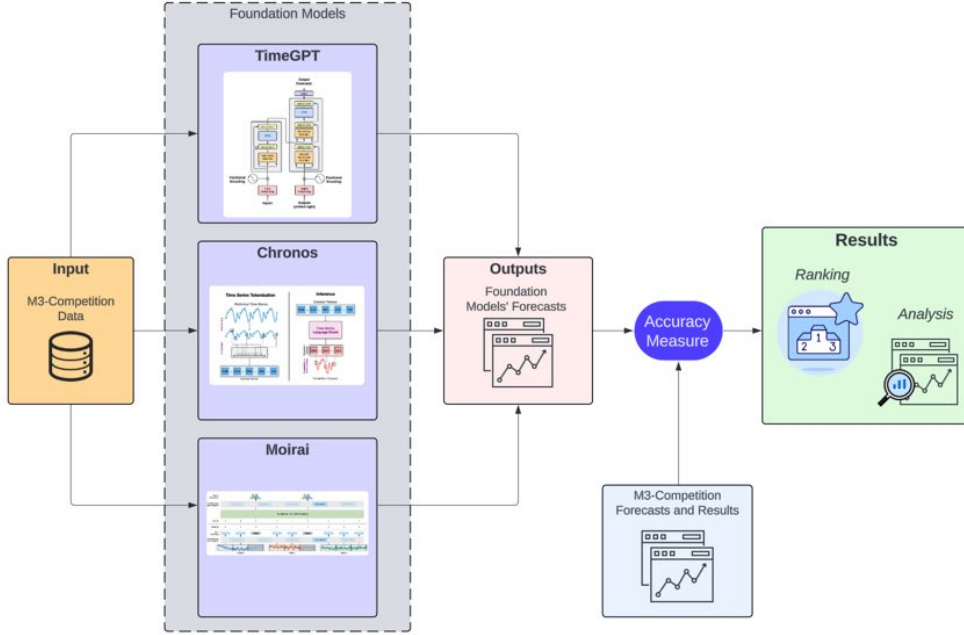


Figure 11: Illustration of the methodology used in this thesis.

4.2 Competition Selection

In selecting an appropriate time series forecasting competition for zero-shot forecasting, several factors were considered to ensure the chosen competition aligned with our research goals. A primary criterion was the transparency of the competition’s processes. It was essential that the methodology behind the results, particularly how they were derived and evaluated, was openly available and documented. This transparency was fundamental for several reasons. First, having access to both training data and the forecasted outputs from other methods was crucial for a meaningful evaluation. Second,

the documentation of the methods used influences the credibility and trustworthiness of the results. Lastly, comprehensive documentation supports not only the internal validity of our findings but also reflects on the competition’s academic integrity and reputation.

We sought a competition recognized for its academic rigor and respect within the field, attracting participants with a robust knowledge base, which would reinforce the relevance and suitability of our choice.

Moreover, while managing practical considerations, we acknowledged certain constraints and necessities. TimeGPT was supported by a grant of 2,000 USD in API credits from Nixtla. For instance, selecting a competition requiring extensive training data could disadvantage the model if our resources were insufficient to handle the volume. Additionally, the foundation models evaluated in our study require substantial computing power, which increases with the size of the dataset. We also aimed to avoid competitions where unique data engineering techniques or specialized data sources could introduce unfair assessment, such as in the M5-Competition competition, where these factors are significant.

Finally, the data needed to contain diverse time series attributes. Evaluating the generalizability of the foundation models required analyzing their performance across various time frequencies. This approach ensured an assessment of the model’s ability to adapt and forecast effectively over differing time frequencies, providing a robust indication of its versatility and applicability in real-world scenarios.

4.2.1 The M3 Competition

The M3-Competition is part of a series of forecasting competitions spanning back to 1982 when Spyros Makridakis first introduced the competition (International Institute of Forecasters, [n.d.](#)). There have been a total of 5 competitions where different forecasting methods have been compared and evaluated, where the latest, the M5-Competition, took place in 2020 (Makridakis et al., [2019](#)). We decided to evaluate the foundation models in the M3-Competition as this best aligned with our above-mentioned requirements.

The competition uses real-world data from various fields like finance, demographics, and economics. Its academic nature ensures detailed documentation and scholarly discussion, providing a solid basis for evaluation. Importantly, the dataset is simple, including only the target forecast values and their time intervals, with no extra features. It includes different time frequencies—yearly, quarterly, and monthly—giving a broad perspective on zero-shot forecasting. Additionally, the M3-Competition

uses a standard evaluation framework with pre-defined metrics, ensuring consistent and comparable results. These features make the M3-Competition an excellent platform to evaluate our foundation models against established forecasting methods.

Methods Used in the M3-Competition The M3-Competition saw 24 different forecasting methods used. The authors have categorized the methods into six categories, namely Naïve/Simple, Explicit Trend Models, Decomposition, ARIMA/ARARMA model, Expert Systems, and Neural Networks. All categories and methods are shown in Table 3 along with their respective approach to forecasting. Descriptions of the methods and the competitors’ names are found in [Appendix A](#).

Table 3: Forecasting Methods Used in the M3-Competition

Category	Model	Forecasting Approach
Naïve/Simple	NAÏVE2	Random Walk for seasonally adjusted data
	SINGLE*	Single Exponential Smoothing
	HOLT	Holt’s Linear Method
	ROBUST-Trend	Holt’s Linear Method
	WINTER	Holt-Winter’s Seasonal Method
Explicit Trend Models	DAMPEN	Holt’s Linear Method with Dampening
	PP-Autocast	Holt’s Linear Method with Dampening
Decomposition	THETA _{sm}	Hybrid of approaches
	COMB S-H-D	Hybrid of approaches
ARIMA/ARARMA Models	THETA	THETA
	B-J Auto	ARIMA
	AutoBox 1	ARIMA
	AutoBox 2	ARIMA
	AutoBox 3	ARIMA
	AAM 1*	ARIMA
	AAM 2*	ARIMA
	ARARMA	ARARMA
Expert Systems	ForecastPro	Choose between several
	SMARTFCS	Choose between several
	RBF	Choose between several
	Flors-Pearc1	Choose between several
	Flors-Pearc2	Choose between several
	ForcX	Choose between several
Neural Networks	Auto-ANN	Deep Learning

* The models marked with a star are excluded from our analysis. The AAM 1 and AAM 2 models are excluded as they lack forecasts for yearly data. The Single Exponential Smoothing model SINGLE is also excluded due to data quality concerns.

Accuracy Metrics in the M3-Competition Makridakis and Hibon (2000) use five different accuracy metrics to evaluate the models in the M3-Competition. However, due to the disadvantages of

these metrics outlined in section (3.8), they will not be central to our evaluation. Nonetheless, we report results for sMAPE, MdAPE and MdRAE as per Makridakis and Hibon (2000) in Appendix C. The accuracy metrics in the M3 Competition are described in Table 4.

Table 4: Accuracy Metrics Used in the M3-Competition

Metric	Description
sMAPE (Symmetric Mean Absolute Percentage Error)	Averages the absolute percentage error between predicted and actual values.
Average Ranking	Average ranking with sMAPE for each method for each horizon.
Percentage Better	Percentage of times that a given method has smaller forecasting errors than another method.
MdAPE (Median Absolute Percentage Error)	Median of percentage errors between forecasts and actual values.
MdRAE (Median Relative Absolute Error)	Compares median of absolute errors between a forecast and a benchmark forecast (NAïVE2).

4.3 Evaluation

To evaluate the foundation models, we compared their performance against 20 previously assessed methods in the M3-Competition. The evaluation proceeded in two phases. Initially, we created a leaderboard that ranks the models based on their overall performance. This leaderboard provides a clear, comparative overview of how each model performs across all data included in the competition.

In the subsequent phase, we conducted a more detailed analysis of the foundation models’ performance. This deeper examination reviewed their results across different categories of data and time frequencies, as done in the M3-Competition. The aim is not only to discover where foundation models stand in the overall rankings but also to understand their strengths and weaknesses in specific forecasting scenarios compared to the traditional methods. This approach lets us provide both a broad overview and a detailed look at how foundation models perform in the field of zero-shot time series forecasting.

4.3.1 Evaluation Metrics

To assess how effective the foundation models are, we used the evaluation metrics described in section (3.8). We deviated from the methodology of Makridakis and Hibon (2000) by primarily focusing on MASE (Mean Absolute Scaled Error) instead of sMAPE (symmetric Mean Absolute Percentage Error). This choice was informed by the recognition that sMAPE, despite its central role in the original

competition’s evaluation, has acknowledged limitations. Hyndman and Koehler (2006) notably advise against using sMAPE due to these concerns.

However, to provide some consistency with the M3-Competition, we present our results with sMAPE as metric in Table 9 and Appendix C, mirroring the reporting style of Makridakis and Hibon (2000). Here sMAPE is defined in absolute values for the numerator and denominator as done by Makridakis et al. (2019).

4.3.2 Benchmark Model

As relative and scaled errors used in MdRAE and MASE require a model performance as a comparison, we will use the Naïve2 method as a benchmark, consistent with the work of Makridakis and Hibon (2000). The Naïve2 method extends the simplicity of the naïve forecasting model by adjusting for seasonality, providing a baseline that accounts for systematic, seasonal variations in the data. The Naïve2 is defined as:

$$F_{(t+i)} = X_t^* (S_j) \quad (61)$$

The Naïve2 creates a forecast $F_{(t+i)}$ by multiplying a seasonally adjusted time series X_t^* with a seasonal component S_j . The seasonality is removed from the time series as $X_t^* = \frac{X_t}{S_j}$. The seasonal component S_j is computed using classical decomposition method for j periods, which can be quarterly or monthly. This is done for forecast index i , which goes to 6 for yearly data, 8 for quarterly, and 18 for monthly.

4.3.3 Leaderboard

In the M3 Competition, Makridakis and Hibon (2000) chose not to rank the 24 competing forecasting methods. They understood that forecasting is complex, and a single ranking might not accurately show how different methods perform across various categories and time frequencies. This approach highlights the importance of context and specific conditions when evaluating forecasting methods.

Despite these considerations, we believe there is value in providing a straightforward overview of how each method performs. Therefore, we have decided to include a leaderboard in our analysis. This leaderboard will rank the methods based on the median MASE. The intention is not to declare a definitive champion method but rather to offer a look at overall performance, especially regarding the foundation models.

We understand and acknowledge the limitations of a simple ranking system. Subsequent sections of our evaluation explore how each method performs across different data categories and time frequencies. This layered approach provides a broad overview and a detailed evaluation.

4.3.4 Time Frequencies

Given that the M3-Competition includes yearly, quarterly, and monthly data, it is important to consider how the models perform within these frequencies. This analysis helps identify both the strengths and potential shortcomings of the foundation models. To achieve this, we calculated the median MASE for each time frequency subset to determine if there are differences in outcomes. Additionally, we examine how the MASE scores evolve throughout the forecast horizon for these subsets. Finally, we analyze the distributions within each subset to identify any patterns or developments.

4.3.5 Categories

Similar to our analysis of time frequencies, investigating how the models perform across each data category helps us determine if there are specific categories where foundation models excel or fall short. We examined the median MASE scores for each category to evaluate the performance variations. Additionally, we present both mean and median MASE scores for each time frequency and category to offer a comprehensive overview of performance at the most granular level available in the M3-Competition dataset.

4.3.6 Comparison Against Benchmark

We dedicate a section to better understanding the foundation models' performance compared to the benchmark. As the MASE scores are already scaled to the benchmark performance, we deviate from MASE for simplicity.

First, we calculate the *Percentage Better*, which represents the percentage of times each foundation model achieves a MASE score lower than 1, indicating performance superior to the Naïve2 benchmark. This metric does not consider the magnitude by which models outperform the benchmark, only how frequently they do. Secondly, we draw inspiration from Makridakis and Hibon (2000) to investigate how a selection of models performs across different forecasting horizons using average sMAPE scores against the Naïve2 benchmark.

4.4 Statistical Tests

Statistical testing of the results is fundamental to validate the accuracy and reliability of our findings. To achieve this, we will employ multiple methods. First, we will visualize the results by displaying the distribution of median MASE scores obtained through bootstrap sampling. However, resampling methods do not produce independent values for the evaluation metrics and might lead to underestimating the variance in the performances.

Koning et al. (2005) notes the importance of statistical testing for evaluating forecast accuracy, as relying on descriptive statistics can lead to inaccurate conclusions. We follow a methodology proposed by Rainio et al. (2024) for statistical testing in comparing models. First, we want to investigate if the performance of the three foundation models is significantly different. For this comparison, we test using the *Friedman's test*. The null hypothesis here is that the differences in model performances are statistically significant. This test has been performed in its original notation. Secondly, we want to test if the amount of variance the foundation models produce is similar. We use the *Levene's test* to test the null hypothesis that the variances from the model performances are equal. Additionally, we use the *Wilcoxon Signed-Rank test* to test for significant differences between potentially interesting pairs of models. This has the same objective as the Friedman's test, but it is for assessing between any two models rather than a group of models.

These tests have the advantage of not being sensitive to the choice of evaluation metric and can be used for data that is not normally distributed. We used the `scipy.stats` functions `friedmanchisquare`, `levene`, and `wilcoxon`.

4.5 Implementation - TimeGPT

This section covers how the results for TimeGPT were obtained, meaning how the forecasts were made and what data was given to the model. In addition, it is important to mention that it is not entirely reproducible as TimeGPT requires an *application programming interface* (API) key. Nonetheless, all the data and scripts are uploaded to [GitHub](#), meaning interested parties may perform any complimentary analyses.

4.5.1 Access

Using the model requires an API key that may be obtained by request to Nixtla. We were fortunate enough to receive 2,000 USD in credits for academic purposes.

4.5.2 Reproducibility

Since TimeGPT is proprietary and not publicly available, our study cannot be directly replicated by others, which may impact the perceived robustness of our findings. To mitigate this limitation, we have taken steps to enhance the reproducibility of our work as much as possible. Additionally, we have uploaded all the data files utilized in our study to this repository, including the outputs generated by TimeGPT. We used the first release of the model, and forecasts were produced on 18.04.2024. This approach aims to provide transparency and assist others in understanding and verifying our research processes and results.

4.5.3 API Data Integration

Due to its proprietary nature, TimeGPT is not meant to be hosted locally on personal hardware but is accessed through an API. The model’s complexity and the requirement to keep its weights inaccessible necessitate this remote operation. To interact with TimeGPT, data must be transmitted to the model via the API to generate forecasts.

To manage this process efficiently, data, once prepared, is fed into the model using a for loop. This loop is structured to pace the data input, ensuring it adheres to the API’s rate limits. Each batch of data sent receives forecasts, which are then stored along with the corresponding series identifiers. All the resulting data is saved using Apache Parquet format. Parquet was chosen because of its quick read capabilities and its reliable preservation of data types, which helped maintain uniformity in our data processing workflow.

4.6 Implementation - Chronos & Moirai

This section describes the technical approach employed to generate forecasts using the MOIRAI and Chronos models, focusing on data preparation, model configuration, and forecast generation. It is important to note that although our configuration was comprehensive, simplicity was a key goal. Our intuition is that the main advantage of a foundation model is its ability to generalize without requiring extensive configurations. Therefore, we limited ourselves to essential adjustments, aiming to assess the model’s effectiveness in zero-shot forecasting.

4.6.1 Data Preparation

The forecasting process started with loading the time series data from a preprocessed Parquet file. Each time series is uniquely identified and isolated from the larger dataset to ensure the integrity

and independence of each forecasting instance. Unlike Chronos, Moirai cannot handle data in the commonly used Pandas format directly and requires the data in the format of GluonTS, a Python library for probabilistic time series modeling in Python.

4.6.2 Model Configuration

When configuring both the Moirai and Chronos models, we opted for the "base" model size. This choice provided a balance between computational efficiency and predictive performance suitable for our hardware's constraints. While the larger variants of these models could have provided increased accuracy, their demands on computational resources and memory were too intensive for our hardware. Consistent with our goal of simplicity, we aimed to minimize complexity in setting up the models. Our configurations closely followed the default settings recommended in their respective GitHub repositories, ensuring we could assess each model's zero-shot forecasting abilities without unnecessary modifications.

4.6.3 Forecast Generation and Storage

Forecasts were generated by iterating through each unique series in the dataset. For each series, the model considers only the first $N - NF$ records for training, where N is the total number of observations in the series, and NF is the number designated for forecasting. This approach mirrors the training-test split commonly used in time series analysis but is adapted here for the M3-Competition.

Both models output a set of probabilistic distributions for future values. These distributions were processed to extract the median forecast at each time step, giving a central tendency measure that is robust to outliers and extreme values. This was done because we focus on point forecasts rather than on distributions.

The forecasts generated for each series were compiled into a single DataFrame. This DataFrame was systematically updated with each batch of forecasts while maintaining a record of all predictions made during the session. Once all forecasts were generated and recorded, the DataFrame was saved as a Parquet file for further analysis.

5 Data

5.1 Data Collection

The dataset used in our thesis was obtained from the International Institute of Forecasters' official website (International Institute of Forecasters, [n.d.](#)). This dataset encompasses the actual data and the results of the M3-Competition, detailed in (Makridakis & Hibon, [2000](#)).

5.2 Data Quality

The dataset was subjected to rigorous quality checks to identify any potential issues. The data was reliable during our exploratory analysis and cleaning process, with no significant quality concerns. However, we did encounter an inconsistency concerning the dataset's values: it was noted in Makridakis and Hibon ([2000](#)) that all values were supposed to be non-negative, aligning with the requirements of using sMAPE as an accuracy metric. Despite this, we discovered one instance where a value did not meet this criterion. We decided to retain this inconsistency in the dataset to maintain consistency with the M3-Competition's methodology and ensure that our results could be comparable with those from the competition.

5.3 Data Description

The dataset used in this thesis consists of real-world data from various domains, though specific details about the data's origins are not provided. This dataset includes time series data that varies in length depending on its time frequency: yearly series require at least 14 observations, quarterly series need at least 16, and monthly series must have a minimum of 48 observations as per Makridakis and Hibon ([2000](#)). As previously discussed, a desirable feature of our purpose was the exclusion of exogenous variables, a requirement fulfilled by this dataset.

Time series with intervals labeled as "other" are not included in our analysis. This exclusion is necessary because our foundation models require a structured timestamp to produce forecasts. Including these series could lead to unreliable forecasting results, as the results would be more dependent on the frequency we assigned, affecting their validity.

One frequently encountered challenge in time series is managing outliers, which our dataset contains. This allows us to evaluate how well the foundation models handle real-world data scenarios. The boxplots ([12](#)) show that variations do exist while the overall distribution is fairly consistent across different categories.

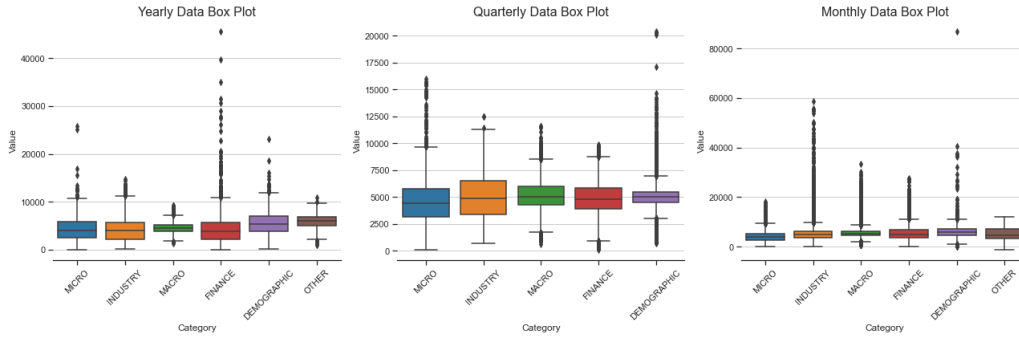


Figure 12: Boxplots of categories across the different time frequencies.

Specifically, some categories are more susceptible to outliers and demonstrate a higher level of skewness in their data distribution. The categories "Demographic" and "Finance" stand out as having more outliers than the other categories.

The table of summary statistics (5) for each category, broken down by frequency (monthly, quarterly, and yearly), reveals interesting aspects of our data. For example, some categories, particularly within the monthly frequency, show significant variance between their minimum and maximum values, indicative of potential outliers. The skewness, inferred from the spread between the 25th and 75th quantiles, suggests a deviation from a normal distribution in several categories, which is affirmed by histogram plots in [Appendix B](#).

Table 5: Summary Statistics of the M3-Competition Dataset

Frequency	Category	Min	Lower Quartile	Mean	Upper Quartile	Max
Monthly	DEMOGRAPHIC	120.00	4,706.00	5,971.10	7,257.25	86,730.00
	FINANCE	10.00	3,643.65	5,265.75	6,663.05	27,505.00
	INDUSTRY	90.00	3,600.00	4,966.40	6,050.00	58,676.00
	MACRO	635.00	4,545.00	5,568.18	6,212.38	33,350.00
	MICRO	100.00	2,600.00	4,082.09	5,260.00	18,100.00
	OTHER	-1,200.00	3,160.00	5,149.64	7,001.90	11,855.20
Quarterly	DEMOGRAPHIC	720.00	4,524.12	5,114.14	5,507.55	20,375.00
	FINANCE	121.00	3,875.00	4,871.71	5,820.82	9,903.33
	INDUSTRY	680.00	3,350.00	5,055.66	6,541.59	12,465.00
	MACRO	650.00	4,267.68	5,178.98	5,976.80	11,601.60
	MICRO	126.00	3,141.00	4,589.06	5,764.40	15,973.00
Yearly	DEMOGRAPHIC	170.80	3,832.00	5,456.84	7,052.70	23,103.30
	FINANCE	30.00	2,175.90	4,352.19	5,642.04	45,525.66
	INDUSTRY	83.10	2,144.75	4,118.87	5,764.00	14,710.40
	MACRO	1,334.00	3,874.75	4,595.86	5,226.50	9,268.50
	MICRO	48.00	2,432.24	4,304.06	5,855.75	25,805.00
	OTHER	1,000.00	4,947.50	5,835.10	6,920.00	10,900.00

As mentioned in (3.1), stationarity concerns the trend and seasonality of the data. To address this,

Augmented Dickey Fuller (ADF) tests were performed for each time series. ADF tests have the null hypothesis that the series are non-stationary (Hyndman & Athanasopoulos, 2021). In our case, this is true for 2,341 of the 2829 unique time series. This high incidence of non-stationarity further confirms that it is representable as real-world data.

5.4 Data Preparation

The M3-Competition dataset was supplied in an Excel file that included actual values as well as forecasts from various models. It had to be adapted to get the data ready for analysis with foundation models. First, we reshaped the dataset from a wide to a long format, which made working with DateTime values and merging operations more straightforward. During this transformation, careful attention was paid to the management of identifier columns to ensure the accuracy and consistency of the data.

One of the benefits of the M3-Competition dataset is that it does not contain exogenous variables, so there was no need for the extra steps of feature engineering or selection. This streamlined the preprocessing work. For greater efficiency in our analysis, we converted all data files to the Apache Parquet format.

5.5 Data Limitations

In the foundation models we evaluate, there is a difference in the time frequency of the data they are trained on. For instance, Moirai training utilized the LOTSA dataset, in which a mere 0.003% represents yearly data. For the models we utilize, the bulk of the training data for these models, except for TimeGPT, where the training data is proprietary, consists mainly of hourly data. Although these models are architected to perform zero-shot forecasting across various frequencies, the amount of hourly data in their training raises questions about their effectiveness at different frequencies. This aspect requires closer examination in the results section, especially regarding its impact on model performance with yearly data. Moreover, the M3 dataset does not extend to time frequencies lower than monthly, so it does not thoroughly test the capabilities to handle the time frequency most represented in their training data.

6 Results

In this section, we present the results of our investigation into the performance of the three foundation models in the M3-Competition. Our research focuses on these models’ performance in zero-shot forecasting, aiming to highlight their strengths, limitations, and comparative performance for the dataset.

6.1 Competition Results

Before evaluating model performance, we provide an overview of the median MASE scores obtained for the foundation models in our evaluation. The median MASE scores show a contrast in performance by the foundation models. THETA is, as mentioned in the original competition, the model to beat for the foundation models to be crowned as superior models. While Chronos is showing a promising result and emerges as a potential contender, TimeGPT and Moirai fall far behind, displaying accuracy levels comparable to the Naive2 benchmark of MASE equal to one.

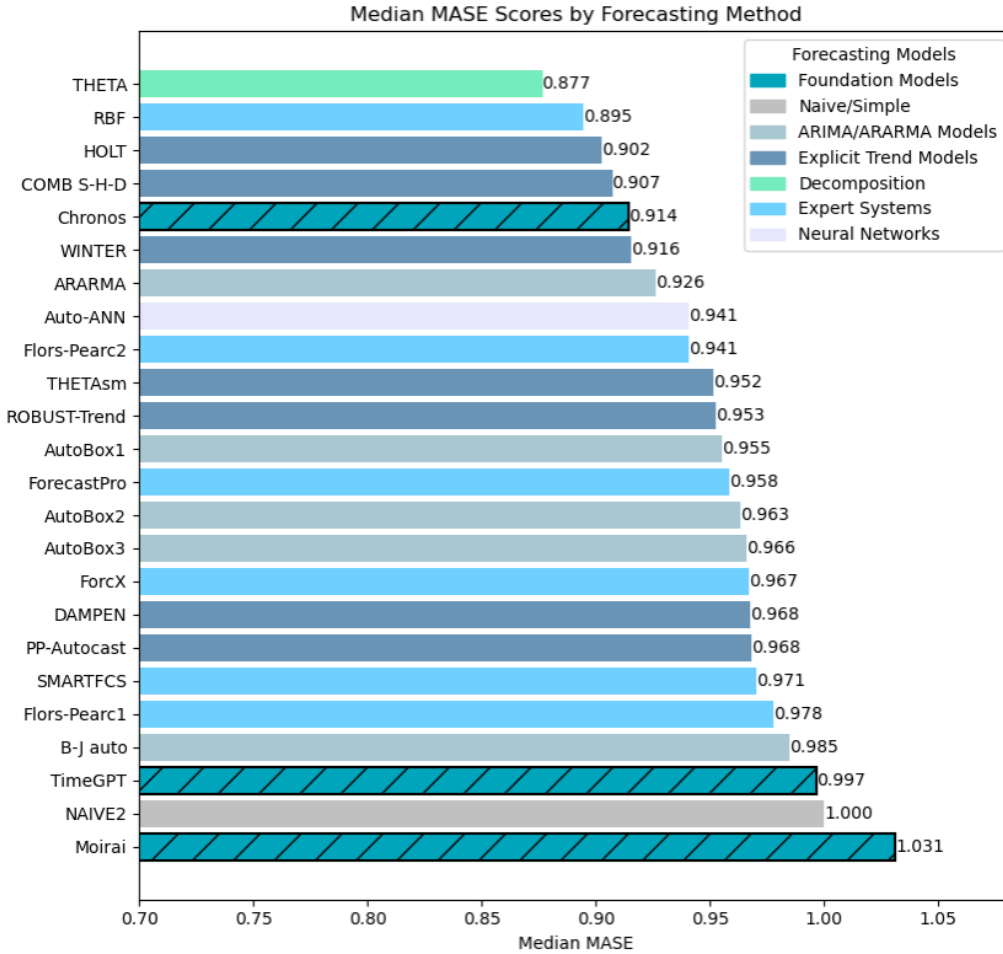


Figure 13: Median MASE scores for all models in the M3-Competition.

The notched boxplot in (14) shows the distribution of the MASE scores for a selection of competitors

in our evaluation with the median as a red line. The selection consists of the top-performing models in the M3-Competition and the foundation models. This showcases that the greater part of the interquartile range (IQR) for TimeGPT and Moirai falls over the benchmark of a MASE score equal to one.




Figure 14: Boxplot of MASE scores for selected models in the M3-Competition.

Note: The outliers for the MASE scores are excluded as all the models have significant outliers in MASE scores which makes the plot unreadable.

6.1.1 Leaderboard Position

As mentioned, the M3-Competition do not provide a leaderboard for its participants in the competition. We, however, present our ranking of the models here to provide a clear overview of model performance. Our findings rank the models as shown in Table (6), with THETA emerging as the top model.

Chronos is the top performer among the foundation models, placing fifth in our evaluation based on the median MASE score. TimeGPT and Moirai present the poorest performance across all evaluated models. It is worth noting that if we were to base our ranking on mean MASE scores, TimeGPT and Moirai would rank slightly higher, and Chronos would place fourth. However, the mean is influenced by outliers in MASE scores, which differ considerably between the models.

6.2 Evaluation

In this section, we build upon the competition results presented in the previous section to conduct a more comprehensive evaluation of the performance of the foundation models. Our evaluation aims to

Table 6: MASE results in the M3-Competition

Rank	Method	Category	Median Score*	Mean Score	Lower Quartile	Upper Quartile
1	THETA	Decomposition	0.88	1.02	0.63	1.10
2	RBF	Expert Systems	0.89	1.12	0.60	1.23
3	HOLT	Explicit Trend	0.90	1.26	0.59	1.28
4	COMB S-H-D	Explicit Trend	0.91	1.01	0.68	1.08
5	Chronos	Foundation Model	0.91	1.05	0.66	1.18
6	WINTER	Explicit Trend	0.92	1.26	0.63	1.23
7	ARARMA	ARIMA/ARARMA	0.93	1.26	0.59	1.36
8	Auto-ANN	Neural Networks	0.94	1.17	0.67	1.26
9	Flors-Pearc2	Expert Systems	0.94	1.17	0.68	1.24
10	THETAsm	Explicit Trend	0.95	1.01	0.80	1.11
11	ROBUST-Trend	Explicit Trend	0.95	1.16	0.62	1.36
12	AutoBox1	ARIMA/ARARMA	0.96	1.33	0.62	1.37
13	ForecastPro	Expert Systems	0.96	1.10	0.63	1.08
14	AutoBox2	ARIMA/ARARMA	0.96	1.31	0.63	1.23
15	AutoBox3	ARIMA/ARARMA	0.97	1.27	0.62	1.38
16	ForcX	Expert Systems	0.97	1.06	0.67	1.10
17	DAMPEN	Explicit Trend	0.97	1.06	0.71	1.08
18	PP-Autocast	Explicit Trend	0.97	1.11	0.70	1.11
19	SMARTFCS	Expert Systems	0.97	1.19	0.63	1.27
20	Flors-Pearc1	Expert Systems	0.98	1.15	0.71	1.16
21	B-J auto	ARIMA/ARARMA	0.99	1.10	0.71	1.12
22	TimeGPT	Foundation Model	1.00	1.23	0.78	1.34
23	Moirai	Foundation Model	1.03	1.26	0.80	1.37

* Rank Determined by the Median MASE Score

provide deeper insights into the strengths, limitations, and practical implications of using foundation models for zero-shot time series forecasting.

6.2.1 Time Frequencies

The foundation models appear to be sensitive to the time frequency of the data. The figure below shows the median MASE scores for the forecasting methods in each subset of yearly, quarterly, and monthly data.

As seen in Figure (15), the median MASE scores for the models are considerably worse for yearly data, with Chronos falling 13 places in ranking and TimeGPT being the worst-performing model by a 0.27 margin. These MASE scores indicate that Moirai and TimeGPT perform worse than the Naïve2 benchmark within this subset, with Chronos narrowly avoiding the same outcome.

The three foundation models' performance is not convincing for quarterly data. While Chronos exhibits a minimal improvement of 0.01 in the median MASE score, resulting in a modest rise in ranking, TimeGPT demonstrates better performance compared to its yearly counterpart but still falls short of achieving a favorable ranking. Moirai is the only one of the three models that performs worse quarterly than yearly, but only by a narrow margin.

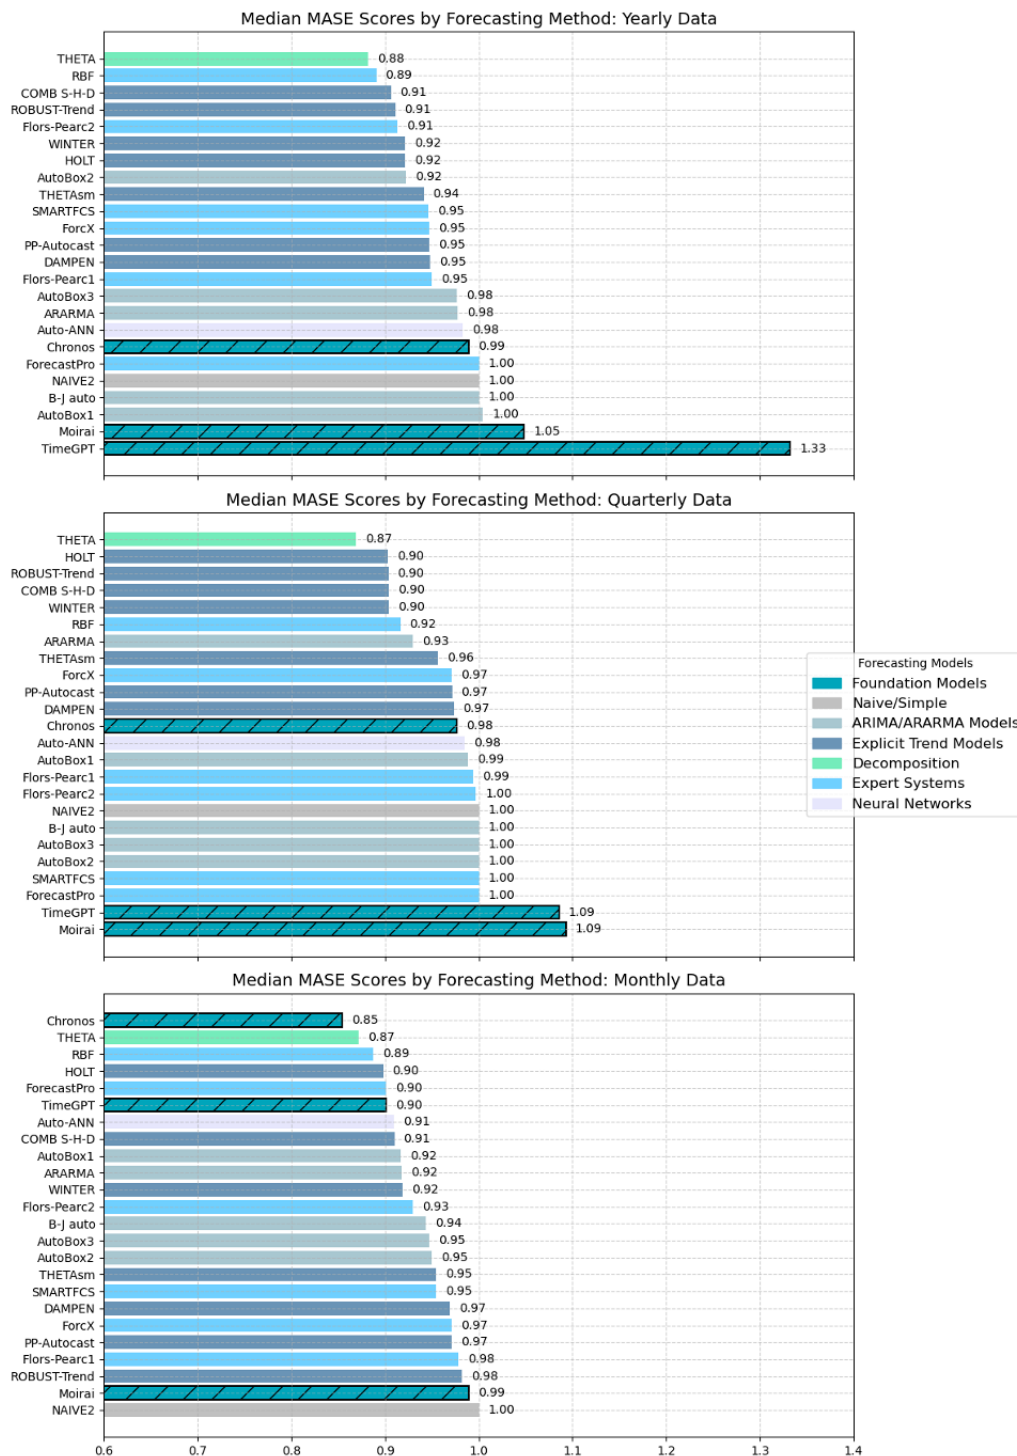


Figure 15: Median MASE scores for different time frequencies.

The median MASE scores changes significantly when analyzing monthly data, with Chronos becoming the top model in this subset. TimeGPT shows a notable improvement, nearly reaching the top five, in contrast to its last-place ranking for yearly data MASE scores. Additionally, [Appendix D](#) shows that the performance of TimeGPT can slightly improve through finetuning for this subset. In contrast, Moirai outperforms the Naive2 benchmark for the first time in our analysis by a small margin of just 0.01.

These observations are underscored by Figure 16, which depicts the variation in the mean MASE scores across the forecast period for the foundation models compared to the Naive2 benchmark. A consistent pattern emerges across all plots: the models underperform the benchmark for yearly data, perform closer to the benchmark for quarterly data, and beats the benchmark for monthly data.




Figure 16: Mean MASE Scores at Each Forecast Step for Different Time Frequencies.

An interesting trend emerges when examining the distributions of MASE scores for the three foundation models across the three time frequencies: the distribution narrows with increasing granularity.

Earlier plots demonstrated how model performance improves as the data frequency shifts from yearly to quarterly to monthly. Figure (17) reveals that, in addition to the improvement in median MASE scores, the scores also become more tightly clustered around the median. The exception is Chronos for monthly data, where, despite the distribution being less concentrated than for quarterly data, the *Interquartile Range* remains well below 1, indicating strong performance for monthly data compared to the benchmark.




Figure 17: Distribution of MASE scores for different time frequencies.

6.2.2 Categories

The performance of the three foundation models varies across different data categories in the M3-Competition. Figure (18) shows the median MASE scores for each model by category. Chronos demonstrates competitiveness with a median MASE score below one in all categories. TimeGPT and Moirai, however, fail to surpass the benchmark in two and four categories, respectively. They each outperform Chronos in only one category, with Moirai doing so by a margin of just 0.03. The most substantial differences between the models are observed in the demographic and macroeconomic categories, where Chronos clearly outperforms both TimeGPT and Moirai, which show similar performance to each other.

As previously demonstrated, the three foundation models are sensitive to the time frequency of the data. It is, therefore, of interest to see the categories for each subset of time frequency to better understand their performance.

Table (7) shows Chronos performing consistently well compared to the other two foundation models, particularly for finance, micro- and macroeconomic data where both the median and mean MASE score place well below the benchmark for monthly data. TimeGPT has a median MASE score below

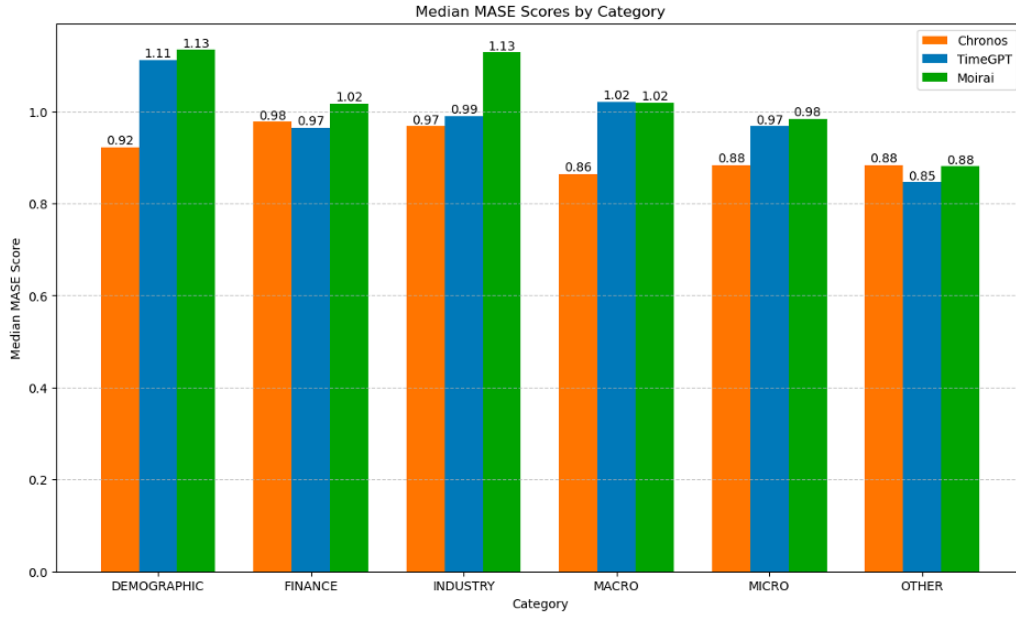


Figure 18: Median MASE Scores for Different Categories.

one for all categories in monthly data and nearly achieves this in mean MASE score, except for industry data. However, it performs remarkably poorly for yearly demographic and macroeconomic data compared to the other models, affecting its ranking in the previously mentioned leaderboard. In the median MASE scores, MOIRAI outperforms TimeGPT in all categories of yearly data. However, this does not hold for other time frequencies, except for quarterly macroeconomic data, where MOIRAI also beats TimeGPT. The category “Other” only contains eleven time series for yearly data and should not hold much weight in this evaluation.

6.2.3 Benchmark

So far, we have used the MASE scores to evaluate the models’ performances. However, as we now directly compare the three foundation models to the benchmark, using MASE might confuse things as the MASE scores are already scaled to the benchmark performance. We use the Percentage Better and the differences in sMAPE scores to simplify the comparisons and provide additional insights. The Percentage Better metric indicates how often the foundation models outperform the benchmark. The differences in sMAPE scores provide a clear percentage-based error comparison between the models and the benchmark.

Table (8) illustrates the Percentage Better across different categories and time series frequencies, giving a view of model performance in varied contexts. We use a weighted average calculation for these percentages to ensure a fair comparison, accounting for varying sample sizes across categories and frequencies. This approach explains how often the foundation models outperform the Naïve2

Table 7: Median and Mean MASE Scores for Different Categories and Frequencies.

Frequency	Category	Chronos		TimeGPT		Moirai	
		median	mean	median	mean	median	mean
Monthly	DEMOGRAPHIC	0.87	1.16	0.95	0.97	1.11	1.21
	FINANCE	0.86	0.92	0.87	0.85	0.98	1.07
	INDUSTRY	0.98	1.05	0.93	1.01	1.11	1.19
	MACRO	0.79	0.98	0.91	0.98	1.03	1.18
	MICRO	0.79	0.82	0.88	0.88	0.92	0.93
	OTHER	0.8	0.84	0.83	0.85	0.88	1.02
Quarterly	DEMOGRAPHIC	1.06	1.11	0.99	1.26	1.05	1.41
	FINANCE	1.02	1.12	1.01	1.21	1.05	1.1
	INDUSTRY	0.96	1.06	1.13	1.45	1.28	1.56
	MACRO	0.95	1.25	1.11	1.36	1.04	1.33
	MICRO	0.97	1.09	1.09	1.31	1.15	1.38
Yearly	DEMOGRAPHIC	0.93	1.14	1.51	2.2	1.18	1.8
	FINANCE	1.06	1.28	1.13	1.32	1.06	1.49
	INDUSTRY	0.95	1.08	1.15	1.43	1.1	1.87
	MACRO	1.02	0.95	1.73	1.84	0.95	1.07
	MICRO	1.01	1.2	1.29	1.45	1.04	1.27
	OTHER	0.99	0.93	0.99	0.98	0.89	0.87

benchmark.

Table (9) shows a comparison of the sMAPE scores between a selection of models against Naive2 as the benchmark with averages for different time horizons (full comparison in [Appendix C](#)). Table (9) underscores the competitiveness of Chronos. Equally, it reveals the shortcomings of TimeGPT and Moirai as contenders despite their improving performance with longer forecast horizons. The forecast horizon, inherently tied to the data frequency, reveals that the foundation models demonstrate progressively better performance as granularity increases. This supports the earlier finding that the foundation models’ performance improves for monthly data.

6.3 Robustness of Results

To visualize how the performance of the three foundation models would have looked under different variations of the dataset from the M3-Competition, we draw 1,000 samples with a sample size of 500 in Figure (19). This further validates our findings, showing that if we were to evaluate the models on random data samples instead of the entire datasets, Chronos would be the preferred model in most cases.

Table 8: Percentage Better than Naïve2 for Different Time Frequencies and Categories.

Frequency	Category	#Time Series	Percentage Better		
			Chronos	TimeGPT	Moirai
Monthly	DEMOGRAPHIC	111	57.66	61.26	44.14
	FINANCE	145	62.07	63.45	53.79
	INDUSTRY	334	53.59	59.88	40.72
	MACRO	312	69.55	64.42	46.79
	MICRO	473	75.48	69.13	63.21
	OTHER	52	67.31	67.31	57.69
Quarterly	DEMOGRAPHIC	57	45.61	50.88	45.61
	FINANCE	76	46.05	48.68	43.42
	INDUSTRY	83	55.42	38.55	26.51
	MACRO	336	53.57	41.96	46.43
	MICRO	204	53.92	39.71	34.8
Yearly	DEMOGRAPHIC	245	55.51	31.84	39.18
	FINANCE	58	37.93	34.48	41.38
	INDUSTRY	102	57.84	34.31	43.14
	MACRO	83	46.99	7.23	60.24
	MICRO	147	46.26	26.53	46.94
	OTHER	11	63.64	54.55	90.91
Weighted Average			59.03	50.44	47.33

Table 9: Differences in sMAPE Scores Against Benchmark.

Method	1	Avg 1-4	Avg 1-6	Avg 1-12	Avg 1-18
THETA	2.15	2.23	2.09	2.37	2.54
ForecastPro	1.93	2.02	1.85	2.18	2.40
Chronos	1.92	1.87	1.54	2.10	2.41
ForcX	1.87	1.86	1.86	2.10	2.01
DAMPEN	1.77	1.58	1.50	1.90	1.93
COMB S-H-D	1.66	1.56	1.54	1.97	2.10
RBF	0.67	1.13	1.35	1.57	1.91
TimeGPT	0.28	-0.08	-0.61	0.92	1.19
Moirai	-0.86	-0.44	-0.25	0.39	0.65




Figure 19: Bootstrap Sampling Distribution of Median MASE Scores

The Friedman’s test reveals a statistically significant difference between the performance of the foundation models. The Wilcoxon Signed-rank test further confirms this is true for any combination of the three models. In addition, the Levene’s test shows that the variances of the three foundation models are unequal. The Wilcoxon Signed-rank test also discovers that all three foundation models perform significantly worse than Theta, the champion model in our evaluation. All statistical tests are done with a 5% confidence interval. More details in [Appendix E](#).

We conclude this section by investigating how the foundation models rank in MASE scores for each time series to see if the placement in the earlier leaderboard is justifiable. Figure (20) shows the count of times each foundation model occupies each rank across all the time series.

Based on MASE rankings, TimeGPT and Moirai tend to place last more frequently than any other position, with the second-to-last position being their next most common ranking. This supports their low leaderboard rankings. TimeGPT ranks first 168 times, which equals 5.9% of the time, making it the third most frequent position for this model. This indicates that TimeGPT can be useful in some cases compared to the other models despite their ranking. Moirai displays a clear trend of increasing counts for worse ranks, with the last three positions being the most frequent. Chronos ranks first most frequently, 6.4% of the time, with second place as the second most frequent position. In fact, Chronos place more often in the top three than any other model in our selection. This justifies its higher leaderboard position.

The plot shows that rankings vary significantly across time series for all three models. This suggests that the leaderboard should not be taken at face value, as the foundation models’ performance varies significantly across the time series. For more details and additional figures, see [Appendix E](#).




Figure 20: Stacked Barchart of *Leaderboard* Placements of Foundation Models.

7 Discussion

7.1 Discussion of Results

The evaluation of the three foundation models—Chronos, TimeGPT, and Moirai—within the context of the M3-Competition demonstrates insights into their performance and characteristics. Overall, the models did not exceed expectations uniformly. Chronos achieved an impressive fifth place, whereas TimeGPT and Moirai placed in the last two positions. This indicates that while foundation models show promise, there is a large discrepancy in their performances.

One notable observation is the sensitivity of model performance to the time frequency of the data. All three models performed better with monthly data than quarterly and yearly data, with Chronos emerging as the champion model for monthly data. This trend may be attributed to the models being trained on data with more granular datasets. Furthermore, the distribution of MASE scores for the models narrows with increasing data granularity, indicating that the models perform better with higher frequency data. When comparing the models across different data categories, Chronos consistently outperformed the Naïve2 benchmark, demonstrating its robustness across varied types of data. In contrast, TimeGPT and Moirai performed worse and more varied, aligning with their performance in the different frequencies of data.

Among the three models, Chronos outperformed the others, demonstrating competitive performance in the M3-Competition. TimeGPT showed notable performance with monthly data but lacked the

same competitiveness for yearly and quarterly data. Moirai, while the least competitive overall, still managed to exceed the benchmark nearly half the time, suggesting that even the low-performing model has potential value in certain contexts.

This evaluation highlights the strengths and weaknesses of the foundation models, emphasizing their varying performances across different data frequencies. While Chronos stands out as the most competitive model, Section (6.3) reveals that our initial leaderboard should not be taken at face value, as both TimeGPT and MOIRAI perform well for certain time series.

7.2 Limitations

Despite the insights provided by our research in this thesis, several limitations should be acknowledged when judging our findings.

First, this thesis exclusively investigates how the foundation models perform in zero-shot forecasting. We do not fine-tune models or use exogenous variables, which could enhance the performance of what we present in our main results. Fine-tuning is a critical step when applying models in real-world scenarios, and its exclusions may limit the generalizability of our findings and do not accurately describe the models' optimal performance. We do, however, present our experiment with fine-tuning TimeGPT in [Appendix D](#), but this is for a subset of the data, and it would therefore be dishonest to give it any weight in our evaluation.

Second, TimeGPT is not an open-source model, which limits transparency and reproducibility. The two other models in our evaluation, Chronos and Moirai, are relatively new and lack extensive documentation. This also limits the knowledge of the models despite our efforts to describe them to the best of our capabilities. It also caps the depth in understanding the models' mechanisms and performance characteristics.

Third, we utilize the M3-Competition dataset, which is reputable and relatively old. Therefore, the data does not necessarily represent current trends and reflect today's most relevant data sources. Also, the dataset is not particularly large. We deem the size to be big enough for this evaluation, but its size limits the ability to generalize our findings to more granular frequencies, such as daily or hourly data and complex streaming data.

Lastly, we do not compare the foundation models to the performance of other advanced and sophisticated forecasting methods, such as gradient boosting or random forest models. While these methods have become widely popular, we have deemed this comparison unnecessary for the scope of this thesis, especially as the Theta model is included, which performed very well on the dataset. Nonethe-

less, the absence of these comparisons can be a limitation in demonstrating the relative performance of foundation models to other new and complex models.

7.3 Future Research

While this thesis helps to better understand the performance of foundation models in zero-shot forecasting, future research on foundation models for time series forecasting could benefit from several different explorations. First, it would be interesting to replicate this study on a larger dataset or competition. Analyzing a more extensive and diverse dataset would help determine whether the observed results hold true across a broader spectrum of data and scenarios. Specifically, using weekly, daily, and even hourly data could provide valuable insights into whether model performance continues to improve with higher frequency data, as suggested by our findings with monthly data.

Additionally, exploring the adaptation and evolution of foundation models and other AI models in the coming years will be crucial. Staying current on upcoming developments and assessing their implications for time series forecasting will be important. New models and techniques are constantly emerging, and it would be of significant interest to the scientific community and practitioners to research what technologies are being adopted, how easily they can be integrated, and what factors contribute to breakthroughs in this field.

It would be valuable to compare these foundation models with other innovative approaches, such as the recently introduced xLSTM modifications to LSTM by Beck et al. (2024). New developments in model architecture are emerging for both Large Language Models and Computer Vision applications. It will be interesting to examine how well these architectures apply to time series, especially considering how well Chronos performed in this study despite minimal efforts to modify the architecture for time series data. This comparison could provide deeper insights into the potential advantages and limitations of different model architectures in handling time series data.

Another promising area for future research is the application of explainable AI (XAI) techniques to foundation models in time series forecasting. Understanding the decision-making processes of these complex models is essential for building trust and understanding their utility in real-world applications. Advancements in large language models (LLMs) are ongoing, as demonstrated by recent work from Templeton et al. (2024), which shows that interpretable features can be extracted from *Claude 3*, a popular LLM. Researchers can strengthen their applicability and reliability in various forecasting scenarios by developing methods to make the predictions of foundation models more transparent and interpretable.

In summary, future research should focus on expanding the scope of data and competitions to validate current findings, continuously adapting to advancements in AI, and emphasizing interpretability to ensure that foundation models are effective and trustworthy in time series forecasting. These directions will help pave the way for more robust, adaptable, and transparent forecasting models.

8 Conclusion

This thesis evaluates the zero-shot forecasting performance of foundation models using the M3-Competition dataset. Utilizing the MASE metric, we compare these models against the Naïve2 benchmark, a simpler forecasting method, to determine if the foundation models offer superior performance.

The analysis reveals that none of the three foundation models outperform the original competition’s champion, the Theta model. However, the foundation model Chronos stands out, ranking fifth overall and surpassing the benchmark, indicating significant promise, especially as an open-source option. In contrast, the other two models, TimeGPT and Moirai, performed worse, ranking at the bottom of our initial ranking. This is particularly interesting because Chronos employs a minimalistic approach to time series forecasting, unlike the other two models. It makes only minor adaptations to a large language model and is the only model trained on synthetic data.

The frequency of the data strongly impacts the models’ performance. The foundation models performed similarly to or worse than the benchmark for yearly and quarterly data, indicating their unsuitability for these data types. Conversely, the models, particularly Chronos and TimeGPT, performed better for monthly data, with Chronos emerging as the best performer among all models. The performance of these models was also more stable for monthly data, suggesting a potential use case for foundation models in this context.

Using foundation models for zero-shot forecasting appears to be a viable approach, though model choice is important. Performance varies by data type, with established methods preferred for yearly and quarterly data. However, foundation models such as Chronos and TimeGPT show strong performance for monthly data, with Chronos notably outperforming all models within this subset. Given its outstanding performance, practitioners should be encouraged to consider Chronos as an option for monthly data forecasting.

References

- Alexis Cook, D., & Holbrook, R. (2021). *Store sales - time series forecasting*. <https://kaggle.com/competitions/store-sales-time-series-forecasting>
- Amazon Web Services. (n.d.). *What is rnn?* <https://aws.amazon.com/what-is/recurrent-neural-network/>
- Ansari, A. F., Stella, L., Turkmen, C., Zhang, X., Mercado, P., Shen, H., Shchur, O., Rangapuram, S. S., Pineda Arango, S., Kapoor, S., Zschiegner, J., Maddix, D. C., Mahoney, M. W., Torkkola, K., Gordon Wilson, A., Bohlke-Schneider, M., & Wang, Y. (2024). Chronos: Learning the language of time series. *arXiv arXiv:2403.07815*. <https://doi.org/10.48550/arXiv.2403.07815>
- Assimakopoulos, V., & Nikolopoulos, K. (2000). The theta model: A decomposition approach to forecasting. *International Journal of Forecasting*, 16(4), 521–530. [https://doi.org/10.1016/S0169-2070\(00\)00066-2](https://doi.org/10.1016/S0169-2070(00)00066-2)
- Ba, J. L., Kiros, J. R., & Hinton, G. E. (2016). Layer normalization. *arXiv arXiv:1607.06450*. <https://doi.org/10.48550/arXiv.1607.06450>
- Bahdanau, D., Cho, K., & Bengio, Y. (2016). Neural machine translation by jointly learning to align and translate. *arXiv arXiv:1409.0473v7*. <https://doi.org/10.48550/arXiv.1409.0473>
- Bai, S., Kolter, J., & Koltun, V. (2018). An empirical evaluation of generic convolutional and recurrent networks for sequence modeling. *arXiv arXiv:1803.01271*. <https://doi.org/10.48550/arXiv.1803.01271>
- Banachewicz, K., Massaron, L., & Goldbloom, A. (2022). *The kaggle book: Data analysis and machine learning for competitive data science*. Packt Publishing.
- Beck, M., Pöppel, K., Spanring, M., Auer, A., Prudnikova, O., Kopp, M., Klambauer, G., Brandstetter, J., & Hochreiter, S. (2024). Xlstm: Extended long short-term memory. *arXiv arXiv:2405.04517*. <https://doi.org/10.48550/arXiv.2405.04517>

- Bommasani, R., Hudson, D. A., Adeli, E., Altman, R., Arora, S., von Arx, S., Bernstein, M. S., Bohg, J., Bosselut, A., Brunskill, E., Brynjolfsson, E., Buch, S., Card, D., Castellon, R., Chatterji, N., Chen, A., Creel, K., Davis, J. Q., Demszky, D., . . . Liang, P. (2022). On the opportunities and risks of foundation models. *arXiv arXiv:2108.07258*. <https://doi.org/10.48550/arXiv.2108.07258>
- Box, G., & Jenkins, G. (1970). *Time series analysis: Forecasting and control*. Holden-Day.
- Brown, R. G. (1959). *Statistical forecasting for inventory control*. McGraw-Hill.
- Brown, R. G. (1963). *Smoothing, forecasting and prediction of discrete time series*. Prentice-Hall.
- Chatfield, C. (2000). *Time-series forecasting*. CRC Press.
- Chen, T., & Guestrin, C. (2016). Xgboost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. <https://doi.org/10.1145/2939672.2939785>
- Cheng, J., Dong, L., & Lapata, M. (2016). Long short-term memory-networks for machine reading. In J. Su, K. Duh, & X. Carreras (Eds.), *Proceedings of the 2016 conference on empirical methods in natural language processing* (pp. 551–561). Association for Computational Linguistics. <https://doi.org/10.18653/v1/D16-1053>
- Davydenko, A., & Fildes, R. (2016). Forecast error measures: Critical review and practical recommendations. <https://doi.org/10.13140/RG.2.1.4539.5281>
- De Gooijer, J. G., & Hyndman, R. J. (2006). 25 years of time series forecasting. *International Journal of Forecasting*, 22(3), 443–473. <https://doi.org/10.1016/j.ijforecast.2006.01.001>
- Deininger, L., Stimpel, B., Yuce, A., Abbasi-Sureshjani, S., Schönenberger, S., Ocampo, P., Korski, K., & Gaire, F. (2022). A comparative study between vision transformers and cnns in digital pathology. <https://doi.org/10.48550/arXiv.2206.00389>

- Deng, J., Dong, W., Socher, R., Li, L.-J., Li, K., & Fei-Fei, L. (2009). Imagenet: A large-scale hierarchical image database. *2009 IEEE Conference on Computer Vision and Pattern Recognition*, 248–255. <https://doi.org/10.1109/CVPR.2009.5206848>
- Dräger, S., & Dunkelau, J. (2022). Evaluating the impact of loss function variation in deep learning for classification. *arXiv arXiv:2210.16003*. <https://doi.org/10.48550/arXiv.2210.16003>
- Freund, Y., & Schapire, R. E. (1997). A decision-theoretic generalization of on-line learning and an application to boosting. *Journal of Computer and System Sciences*, 55(1), 119–139. <https://doi.org/10.1006/jcss.1997.1504>
- Friedman, J. H. (2001). Greedy function approximation: A gradient boosting machine. *The Annals of Statistics*, 29(5), 1189–1232. <https://doi.org/10.1214/aos/1013203451>
- Gardner, E. S., & McKenzie, E. (1985). Forecasting trends in time series. *Management Science*, 31(10), 1237–1246. <https://doi.org/10.1287/mnsc.31.10.1237>
- Garza, A., & Mergenthaler-Canseco, M. (2023). Timegpt-1. <https://doi.org/10.48550/arXiv.2310.03589>
- Gers, F., Schmidhuber, J., & Cummins, F. (2000). Learning to forget: Continual prediction with lstm. *Neural computation*, 12, 2451–71. <https://doi.org/10.1162/089976600300015015>
- Glorot, X., & Bengio, Y. (2010). Understanding the difficulty of training deep feedforward neural networks. In Y. W. Teh & M. Titterton (Eds.), *Proceedings of the thirteenth international conference on artificial intelligence and statistics* (pp. 249–256, Vol. 9). PMLR. <https://proceedings.mlr.press/v9/glorot10a.html>
- Goodwin, P., & Lawton, R. (1999). On the asymmetry of the symmetric mape. *International Journal of Forecasting*, 15(4), 405–408. [https://doi.org/10.1016/S0169-2070\(99\)00007-2](https://doi.org/10.1016/S0169-2070(99)00007-2)
- Green, K. C., & Armstrong, J. S. (2015). Simple versus complex forecasting: The evidence. *Journal of Business Research*, 68(8), 1678–1685. <https://doi.org/10.1016/j.jbusres.2015.03.026>

- He, K., Zhang, X., Ren, S., & Sun, J. (2015). Delving deep into rectifiers: Surpassing human-level performance on imagenet classification. *2015 IEEE International Conference on Computer Vision (ICCV)*, 1026–1034. <https://doi.org/10.1109/ICCV.2015.123>
- Hendrycks, D., & Gimpel, K. (2016). Gaussian error linear units (gelus). *arXiv arXiv:1606.08415v5*. <https://doi.org/10.48550/arXiv.1606.08415>
- Hibon, M., & Makridakis, S. (2000). M3-competition. *International Journal of Forecasting*, 16. [https://doi.org/10.1016/S0169-2070\(00\)00078-9](https://doi.org/10.1016/S0169-2070(00)00078-9)
- Hochreiter, S., & Schmidhuber, J. (1997). Long short-term memory. *Neural computation*, 9, 1735–80. <https://doi.org/10.1162/neco.1997.9.8.1735>
- Holt, C. E. (1957). Forecasting seasonals by exponentially weighted moving averages. <https://doi.org/10.1016/j.ijforecast.2003.09.015>
- Hyndman, R., & Athanasopoulos, G. (2021). *Forecasting: Principles and practice* (3rd). OTexts.
- Hyndman, R., & Koehler, A. B. (2006). Another look at measures of forecast accuracy. *International Journal of Forecasting*, 22(4), 679–688. <https://doi.org/10.1016/j.ijforecast.2006.03.001>
- International Institute of Forecasters. (n.d.). *Time series data*. Retrieved February 9, 2024, from <https://forecasters.org/resources/time-series-data/>
- Ivanov, A., Dryden, N., Ben-Nun, T., Li, S., & Hoefler, T. (2021). Data movement is all you need: A case study on optimizing transformers. *arXiv arXiv:2007.00072v3*. <https://doi.org/10.48550/arXiv.2007.00072>
- Jiang, A. Q., Sablayrolles, A., Mensch, A., Bamford, C., Chaplot, D. S., de las Casas, D., Bressand, F., Lengyel, G., Lample, G., Saulnier, L., Lavaud, L. R., Lachaux, M.-A., Stock, P., Scao, T. L., Lavril, T., Wang, T., Lacroix, T., & Sayed, W. E. (2023). Mistral 7b. *arXiv arXiv:2310.06825*. <https://doi.org/10.48550/arXiv.2310.06825>
- Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., & Liu, T.-Y. (2017). Lightgbm: A highly efficient gradient boosting decision tree. *Advances in Neural Information*

Processing Systems, 30. https://proceedings.neurips.cc/paper_files/paper/2017/file/6449f44a102fde848669bdd9eb6b76fa-Paper.pdf

Kingma, D., & Ba, J. (2014). Adam: A method for stochastic optimization. *International Conference on Learning Representations*. <https://doi.org/10.48550/arXiv.1412.6980>

Koning, A. J., Franses, P. H., Hibon, M., & Stekler, H. O. (2005). The m3 competition: Statistical tests of the results. *International Journal of Forecasting*, 21(3), 397–409. <https://doi.org/10.1016/j.ijforecast.2004.10.003>

Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). Imagenet classification with deep convolutional neural networks. In F. Pereira, C. Burges, L. Bottou, & K. Weinberger (Eds.), *Advances in neural information processing systems* (Vol. 25). Curran Associates, Inc. https://papers.nips.cc/paper_files/paper/2012/file/c399862d3b9d6b76c8436e924a68c45b-Paper.pdf

Lai, G., Chang, W.-C., Yang, Y., & Liu, H. (2018). Modeling long- and short-term temporal patterns with deep neural networks, 95–104. <https://doi.org/10.1145/3209978.3210006>

Lecun, Y., Bottou, L., Bengio, Y., & Haffner, P. (1998). Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11), 2278–2324. <https://doi.org/10.1109/5.726791>

Makridakis, S., & Hibon, M. (2000). The m3-competition: Results, conclusions and implications. *International Journal of Forecasting*, 16(4), 451–476. [https://doi.org/10.1016/S0169-2070\(00\)00057-1](https://doi.org/10.1016/S0169-2070(00)00057-1)

Makridakis, S., Spiliotis, E., & Assimakopoulos, V. (2019). The m4 competition: 100,000 time series and 61 forecasting methods. *International Journal of Forecasting*, 36. <https://doi.org/10.1016/j.ijforecast.2019.04.014>

McCulloch, W. S., & Pitts, W. (1943). A logical calculus of the ideas immanent in nervous activity. *The bulletin of mathematical biophysics*, 5(4), 115–133.

Meade, N. (2000). A note on the Robust Trend and ARARMA methodologies used in the M3 Competition. *International Journal of Forecasting*, 16(4), 517–519. [https://doi.org/10.1016/S0169-2070\(00\)00073-X](https://doi.org/10.1016/S0169-2070(00)00073-X)

- Muth, J. F. (1960). Optimal properties of exponentially weighted forecasts. *Journal of the American Statistical Association*, 55(290), 299–306.
- N., J., & Wold, H. (1939). A study in the analysis of stationary time series. *Journal of the Royal Statistical Society*, 102(2), 295–298.
- Nixtla. (n.d.). *Key concepts*. <https://docs.nixtla.io/docs/key-concepts>
- Pan, S. J., & Yang, Q. (2010). A survey on transfer learning. *IEEE Transactions on Knowledge and Data Engineering*, 22(10), 1345–1359. <https://doi.org/10.1109/TKDE.2009.191>
- Parzen, E. (1982). Ararma models for time series analysis and forecasting. *Journal of Forecasting*, 1(1), 67–82. <https://doi.org/10.1002/for.3980010108>
- Pope, R., Douglas, S., Chowdhery, A., Devlin, J., Bradbury, J., Levskaya, A., Heek, J., Xiao, K., Agrawal, S., & Dean, J. (2022). Efficiently scaling transformer inference. *arXiv arXiv:2211.05102*. <https://doi.org/10.48550/arXiv.2211.05102>
- Raffel, C., Shazeer, N., Roberts, A., Lee, K., Narang, S., Matena, M., Zhou, Y., Li, W., & Liu, P. J. (2020). Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of Machine Learning Research*, 21(140), 1–67. <http://jmlr.org/papers/v21/20-074.html>
- Rainio, O., Teuho, J., & Klén, R. (2024). Evaluation metrics and statistical tests for machine learning. *Scientific Reports*, 14(1), 6086. <https://doi.org/10.1038/s41598-024-56706-x>
- Raschka, S., Liu, Y., Mirjalili, V., & Dzhulgakov, D. (2022). *Machine learning with pytorch and scikit-learn: Develop machine learning and deep learning models with python*. Packt Publishing.
- Slutzky, E. (1937). The summation of random causes as the source of cyclic processes. *Econometrica*, 5(2), 105–146.
- Staudemeyer, R., & Morris, E. (2019). Understanding lstm – a tutorial into long short-term memory recurrent neural networks. *arXiv arXiv:1909.09586*. <https://doi.org/10.48550/arXiv.1909.09586>

- Su, Y., & Kuo, C. (2018). On extended long short-term memory and dependent bidirectional recurrent neural network. *Neurocomputing*, 356. <https://doi.org/10.1016/j.neucom.2019.04.044>
- Templeton, A., Conerly, T., Marcus, J., Lindsey, J., Bricken, T., Chen, B., Pearce, A., Citro, C., Ameisen, E., Jones, A., Cunningham, H., Turner, N. L., McDougall, C., MacDiarmid, M., Freeman, C. D., Sumers, T. R., Rees, E., Batson, J., Jermyn, A., ... Henighan, T. (2024). Scaling monosemanticity: Extracting interpretable features from claude 3 sonnet. *Transformer Circuits Thread*. <https://transformer-circuits.pub/2024/scaling-monosemanticity/index.html>
- Umar Jamil. (2023). *Attention is all you need (transformer) - model explanation (including math), inference and training*. <https://www.youtube.com/watch?v=bCz4OMemCcA>
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30. https://proceedings.neurips.cc/paper_files/paper/2017/file/3f5ee243547dee91fbd053c1c4a845aa-Paper.pdf
- Weiss, K., Khoshgoftaar, T. M., & Wang, D. (2016). A survey of transfer learning. *Journal of Big Data*, 3(1), 9. <https://doi.org/10.1186/s40537-016-0043-6>
- Wen, Q., Zhou, T., Zhang, C., Chen, W., Ma, Z., Yan, J., & Sun, L. (2023). Transformers in time series: A survey. *arXiv arXiv:2202.07125*. <https://doi.org/10.48550/arXiv.2202.07125>
- Werbos, P. (1990). Backpropagation through time: What it does and how to do it. *Proceedings of the IEEE*, 78, 1550–1560. <https://doi.org/10.1109/5.58337>
- Willmott, C. J., & Matsuura, K. (2005). Advantages of the mean absolute error (mae) over the root mean square error (rmse) in assessing average model performance. *Climate Research*, 30(1), 79–82. <https://doi.org/10.3354/cr030079>
- Winters, P. R. (1960). Forecasting sales by exponentially weighted moving averages. *Management Science*, 6, 324–342. <https://doi.org/10.1287/mnsc.6.3.324>

- Woo, G., Liu, C., Kumar, A., Xiong, C., Savarese, S., & Sahoo, D. (2024). Unified training of universal time series forecasting transformers. *arXiv preprint arXiv:2402.02592*. <https://doi.org/10.48550/arXiv.2402.02592>
- Wu, H., Xu, J., Wang, J., & Long, M. (2021). Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting. In M. Ranzato, A. Beygelzimer, Y. Dauphin, P. Liang, & J. W. Vaughan (Eds.), *Advances in neural information processing systems* (pp. 22419–22430, Vol. 34). Curran Associates, Inc. https://proceedings.neurips.cc/paper_files/paper/2021/file/bcc0d400288793e8bdcd7c19a8ac0c2b-Paper.pdf
- Yang, Y., Jiao, L., Liu, X., Liu, F., Yang, S., Feng, Z., & Tang, X. (2022). Transformers meet visual learning understanding: A comprehensive review. *arXiv arXiv:2203.12944*. <https://doi.org/10.48550/arXiv.2203.12944>
- Yule, G. U. (1927). On a method of investigating periodicities in disturbed series, with special reference to wolfer’s sunspot numbers. *Philosophical Transactions of the Royal Society A*, 226, 267–298.
- Zeng, A., Chen, M., Zhang, L., & Xu, Q. (2022). Are transformers effective for time series forecasting? *arXiv arXiv:2205.13504*. <https://doi.org/10.48550/arXiv.2205.13504>
- Zhou, H., Zhang, S., Peng, J., Zhang, S., Li, J., Xiong, H., & Zhang, W. (2020). Informer: Beyond efficient transformer for long sequence time-series forecasting. *arXiv arXiv:2012.07436v3*. <https://doi.org/10.48550/arXiv.2012.07436>

Appendix A

Methods used in the M3-Competition

Table (10) shows an overview of the methods used in the M3 Competition with the short descriptions of Makridakis and Hibon (2000). A longer description of most of the models, those that is not used in earlier M- Competitions, can be found in the appendix of the paper. Further explanation of the Robust-Trend and ARARMA methodologies used is provided by Meade (2000).

Table 10: Description of Forecasting Methods Used in the M3-Competition.

Category	Methods	Competitors	Description
Naïve/ Simple	Naïve 2 Single	M. Hibon M. Hibon	Naïve forecast where seasonality is first removed. Single Exponential Smoothing.
Explicit Trend Models	Holt	M. Hibon	Automatic Holt's Linear Exponential smoothing.
	Robust-Trend	N. Meade	Nonparametric version of Holt's linear model with median based estimate of trend.
	Winter	M. Hibon	Holt-Winter's linear and seasonal exponential smoothing.
	Dampen	M. Hibon	Dampen Trend Exponential Smoothing.
	PP autocast THETA-sm	H. Levenbach V. Assimakopoulos	Dampen Trend Exponential Smoothing. Successive smoothing plus a set of rules for dampening the trend.
	Comb S/H/D	M. Hibon	Combining Single/Holt/Dampen methods.
Decomposition	THETA	V. Assimakopoulos	Specific decomposition technique, projection and combination of the individual components.
ARIMA/ ARARMA Models	BJ-automatic	M. Hibon	Box Jenkins methodology of "Business Forecast System".
	AUTOBOX 1	D. Reilly	Robust ARIMA univariate Box-Jenkins with/without Intervention Detection.
	AUTOBOX 2	D. Reilly	Robust ARIMA univariate Box-Jenkins with/without Intervention Detection.
	AUTOBOX 3	D. Reilly	Robust ARIMA univariate Box-Jenkins with/without Intervention Detection.
	AAM 1	G. Melard & J. M. Pasteels	Automatic ARIMA modeling with/without intervention analysis.
	AAM 2	G. Melard & J. M. Pasteels	Automatic ARIMA modeling with/without intervention analysis.
	ARARMA	N. Meade	Automated Parzen's methodology with Auto regressive filter.
Expert Systems	ForecastPRO	R. Goodrich & E. Stellwagen	Selects from among several methods: Exponential Smoothing/Box Jenkins/Poisson and negative binominal models/Croston's method/Simple Moving Average.
	SMARTFCs	C. Smart	Automatic Forecasting Expert System which conducts a forecasting tournament among 4 Exponential Smoothing and 2 Moving Average methods.
	RBF	M. Adya, S. Armstrong, F. Collopy & M. Kennedy	Rule-Based forecasting: using 3 methods – random walk, linear regression and Holt's to estimate level and trend, involving corrections, simplification, automatic feature identification and recalibration.
	FLORES-PEARCE1	B. Flores & S. Pearce	Expert system that chooses among 4 methods based on the characteristics of the data.
	FLORES-PEARCE2 ForecastX	B. Flores & S. Pearce J. Galt	Expert system that chooses among 4 methods based on the characteristics of the data. Runs tests for seasonality and outliers and selects from among several methods: Exponential Smoothing, Box-Jenkins and Croston's method.
Neural Networks	Automat ANN	K. Ord & S. Balkin	Automated Artificial Neural Networks for forecasting purposes.

Appendix B

Data Distribution By Category

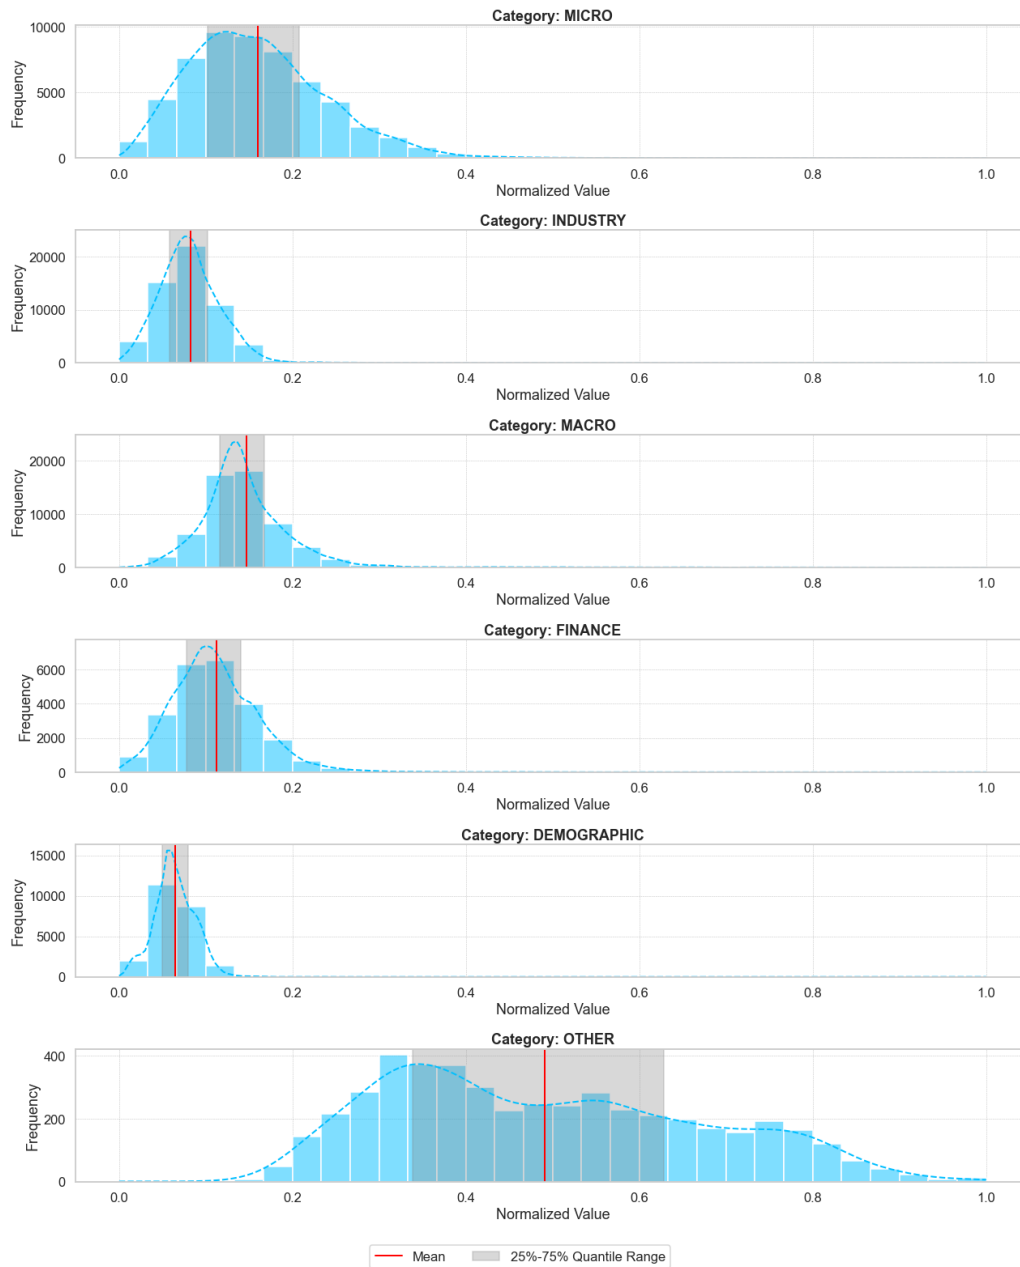


Figure 21: Distribution of the M3-Competition Data by Category.

Appendix C

Additional Results

In this appendix, we provide the scores in other evaluation metrics. This is not intended as a supplementary analysis, but simply offers the reader the opportunity to evaluate the performance of the models in style with the original M3-Competition. We do not rely on findings from these metrics in our evaluation as they can be more sensitive to skewness and outliers. However, the following three tables support our findings with Chronos being in the top three when averaging for all forecast horizons (Avg 1-18) regardless of forecasting metric, while TimeGPT and Moirai falls short.

Table 11: Average Symmetric MAPE with All Data.

Method	Forecast Horizon																		Average of Forecast Horizon						# obs
	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	Avg 1-4	Avg 1-6	Avg 1-8	Avg 1-12	Avg 1-15	Avg 1-18	
NAIVE2	10.96	11.8	14.14	15.6	15.51	16.3	14.95	14.93	15.85	15.95	16.73	15.99	18.09	18.41	19.32	21.29	19.64	20.7	13.13	14.05	14.27	14.89	15.64	16.45	2829
SINGLE	9.9	11.05	13.18	14.59	14.72	15.36	13.74	13.65	13.88	14.0	14.81	14.52	16.19	17.21	18.32	20.18	18.44	19.37	12.18	13.13	13.27	13.62	14.34	15.17	2829
HOLT	9.41	10.93	13.49	15.44	16.16	16.92	14.75	14.93	14.4	15.67	15.4	15.32	17.15	17.85	19.45	20.07	20.07	21.06	12.32	13.72	14.0	14.4	15.15	16.03	2829
DAMPEN	9.19	10.4	12.54	14.08	14.26	14.85	13.22	12.98	13.21	13.34	14.0	13.85	15.63	16.26	17.49	19.43	17.8	18.87	11.55	12.55	12.69	12.99	13.69	14.52	2829
WINTER	9.55	11.0	13.65	15.52	16.21	17.03	14.77	15.06	14.6	15.66	15.5	15.24	17.22	17.9	19.58	20.18	20.01	21.07	12.43	13.83	14.1	14.48	15.23	16.1	2829
COMB S-H-D	9.3	10.45	12.53	14.0	14.16	14.62	13.09	12.9	13.07	13.58	13.69	13.63	15.41	16.19	17.28	18.85	17.36	18.26	11.57	12.51	12.63	12.92	13.59	14.35	2829
B-J auto	9.65	10.86	12.69	14.42	14.51	15.2	13.52	13.44	13.41	13.38	14.52	14.06	16.25	16.94	17.76	19.72	18.15	19.26	11.9	12.89	13.04	13.3	14.04	14.87	2829
AutoBox1	10.3	11.59	13.65	15.77	16.62	17.4	14.86	14.77	14.36	14.33	15.49	15.36	17.53	18.1	19.09	20.42	20.18	20.39	12.83	14.22	14.37	14.54	15.28	16.12	2829
AutoBox2	10.02	10.87	12.76	14.51	14.43	15.47	14.21	13.81	14.63	14.7	15.53	15.18	17.56	17.41	18.37	20.74	19.01	19.99	12.04	13.01	13.26	13.84	14.63	15.51	2829
AutoBox3	10.23	11.82	13.62	15.48	16.79	17.38	15.53	15.43	14.93	15.81	16.12	16.68	18.52	18.32	19.83	21.54	22.6	21.99	12.79	14.22	14.54	14.99	15.77	16.81	2829
ROBUST-Trend	11.03	11.73	13.91	15.42	15.76	16.64	15.75	16.0	18.28	18.69	18.82	18.05	20.93	21.46	23.24	24.1	23.21	25.75	13.02	14.08	14.53	15.84	17.05	18.26	2829
ARARMA	10.19	11.39	13.12	14.81	15.14	16.21	14.53	14.55	14.87	15.01	15.85	15.21	17.51	17.25	18.48	20.01	18.66	20.32	12.38	13.48	13.74	14.24	14.94	15.73	2829
Auto-ANN	9.45	10.98	12.37	14.45	14.46	16.24	13.84	14.0	13.96	14.89	14.71	14.81	17.03	16.83	17.55	19.29	18.25	19.81	11.81	12.99	13.22	13.68	14.37	15.16	2829
Flors-Pearc1	9.61	10.91	13.14	15.04	15.28	15.82	14.7	14.33	14.43	14.29	14.59	14.43	16.68	18.19	19.13	21.22	20.13	21.01	12.18	13.3	13.6	13.88	14.7	15.72	2829
Flors-Pearc2	10.49	11.48	13.32	14.61	14.56	15.2	13.58	13.41	13.26	13.82	14.31	14.37	16.42	17.59	18.23	20.12	18.49	20.24	12.48	13.28	13.33	13.53	14.31	15.19	2829
PP-Autocast	9.56	10.56	12.71	14.21	14.48	15.24	13.82	13.86	14.58	14.29	15.19	14.39	16.47	16.99	17.93	19.78	18.3	19.85	11.76	12.79	13.05	13.57	14.29	15.12	2829
ForecastPro	9.03	9.98	11.89	13.54	13.92	14.85	12.96	13.14	13.02	13.0	13.9	13.31	15.28	15.43	16.43	18.17	16.8	18.29	11.11	12.2	12.41	12.71	13.31	14.05	2829
SMARTFCS	9.59	10.69	12.45	14.05	14.53	15.59	13.67	13.51	13.86	15.19	15.01	14.94	16.55	16.58	18.01	18.43	19.32	19.42	11.7	12.82	13.01	13.59	14.28	15.08	2829
THETA _{asm}	9.91	11.1	13.09	14.34	15.31	16.21	14.51	13.9	14.47	14.81	14.83	14.38	16.7	15.7	17.72	19.26	18.07	19.28	12.11	13.33	13.55	13.9	14.47	15.2	2829
THETA	8.81	9.99	11.77	13.03	13.66	14.5	12.8	12.45	13.17	13.4	13.48	13.23	15.42	15.23	16.36	17.77	16.88	18.36	10.9	11.96	12.13	12.52	13.15	13.91	2829
RBF	10.29	10.95	12.88	13.88	13.62	14.58	13.65	13.19	14.16	14.47	14.14	14.05	16.1	15.78	17.26	18.28	16.76	17.76	12.0	12.7	12.88	13.32	13.93	14.54	2829
ForcX	9.09	10.26	12.09	13.65	13.7	14.37	12.9	13.06	13.06	13.43	13.89	13.97	15.77	16.65	17.85	19.35	18.13	18.79	11.27	12.19	12.39	12.79	13.58	14.44	2829
Chronos	9.04	9.95	12.21	13.84	14.15	15.87	12.75	12.86	12.86	13.19	13.27	13.48	15.37	15.77	15.97	17.63	16.49	18.07	11.26	12.51	12.58	12.79	13.37	14.04	2829
TimeGPT	10.68	11.84	14.04	16.26	16.75	18.41	13.25	13.3	12.86	13.26	13.42	13.6	16.59	16.62	17.3	19.65	17.96	18.91	13.21	14.66	14.32	13.97	14.55	15.26	2829
Moirai	11.82	12.84	14.17	15.46	15.52	15.98	14.62	14.03	14.0	15.1	15.33	15.1	17.38	17.11	18.16	20.11	18.45	19.24	13.57	14.3	14.31	14.5	15.11	15.8	2829

Note: Top Three Scores in Each Column are Bolded.

Table 12: Median Absolute Percentage Error with All Data.

Method	Forecast Horizon																		Average of Forecast Horizon						# obs
	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	Avg 1-4	Avg 1-6	Avg 1-8	Avg 1-12	Avg 1-15	Avg 1-18	
NAIVE2	3.66	4.75	6.05	7.06	7.57	8.47	7.67	7.7	7.82	8.22	8.39	7.79	9.35	9.49	10.24	10.92	10.82	11.62	5.38	6.26	6.62	7.1	7.62	8.2	2829
HOLT	3.35	4.33	5.39	6.59	7.19	7.42	6.6	6.55	6.43	7.46	6.89	7.23	8.26	7.47	8.75	9.13	9.29	10.32	4.92	5.71	5.93	6.29	6.66	7.15	2829
DAMPEN	3.28	4.04	5.1	6.51	6.79	7.42	6.81	6.88	6.73	6.8	6.73	7.47	8.28	8.27	9.69	10.4	9.75	10.51	4.73	5.52	5.85	6.21	6.72	7.3	2829
WINTER	3.42	4.37	5.5	6.51	7.25	7.49	6.66	6.84	6.5	7.23	7.06	7.04	8.52	7.81	9.05	9.26	9.36	10.14	4.95	5.76	6.01	6.32	6.75	7.22	2829
COMB S-H-D	3.35	4.18	5.15	6.47	6.8	7.19	6.58	6.7	6.43	6.84	6.7	6.87	8.26	8.1	9.05	9.39	9.13	9.83	4.79	5.52	5.8	6.1	6.58	7.06	2829
B-J auto	3.29	4.23	5.37	6.6	6.78	7.26	6.56	6.48	6.86	6.94	6.64	7.19	8.24	8.55	9.21	10.44	9.65	10.65	4.87	5.59	5.82	6.18	6.68	7.27	2829
AutoBox1	3.72	4.66	5.75	6.61	7.22	7.93	6.64	7.03	6.45	6.76	7.2	7.04	8.5	8.2	9.17	9.73	9.75	10.55	5.18	5.98	6.2	6.42	6.86	7.38	2829
AutoBox2	3.39	4.41	5.22	6.35	6.43	7.23	6.78	6.36	6.77	6.71	6.81	7.09	8.38	8.57	9.41	9.99	9.4	10.26	4.84	5.5	5.77	6.13	6.66	7.2	2829
AutoBox3	3.62	4.74	5.7	6.8	7.17	7.74	6.87	7.27	6.9	7.03	6.98	7.08	9.06	8.4	8.91	10.05	10.31	10.45	5.22	5.96	6.24	6.49	6.95	7.5	2829
ROBUST-Trend	3.5	4.37	5.24	6.28	7.1	7.55	7.17	7.13	7.67	8.32	8.64	7.47	8.87	8.87	9.72	10.02	10.49	11.68	4.85	5.67	6.04	6.7	7.19	7.78	2829
ARARMA	3.42	4.4	5.32	6.64	7.07	7.83	6.56	7.01	7.16	6.93	6.91	7.08	8.14	8.42	8.85	9.49	9.28	10.18	4.94	5.78	6.03	6.36	6.78	7.26	2829
Auto-ANN	3.4	4.64	5.55	6.82	7.16	7.88	6.89	7.06	7.06	7.35	7.05	7.38	8.64	9.09	9.53	10.27	9.83	11.13	5.1	5.91	6.18	6.52	7.03	7.6	2829
Flors-Pearc1	3.49	4.11	5.13	6.6	6.99	7.31	7.14	7.0	6.64	7.28	6.88	7.46	8.65	8.98	9.8	10.15	10.13	10.73	4.83	5.6	5.97	6.34	6.9	7.47	2829
Flors-Pearc2	4.41	4.96	5.7	7.06	6.96	7.17	6.58	6.94	6.42	7.09	6.74	7.62	8.14	8.61	9.52	9.96	9.72	10.43	5.53	6.04	6.22	6.47	6.93	7.45	2829
PP-Autocast	3.18	4.14	5.37	6.72	7.06	7.6	6.93	6.75	6.92	7.01	7.13	7.37	8.28	8.7	9.84	10.44	10.13	10.98	4.85	5.68	5.97	6.35	6.87	7.48	2829
ForecastPro	3.19	4.09	4.86	6.21	6.54	7.19	6.3	6.53	6.19	6.43	6.63	6.65	7.69	7.69	8.62	9.55	8.97	9.66	4.59	5.35	5.61	5.9	6.32	6.83	2829
SMARTFCS	3.94	4.63	5.42	6.55	6.8	7.26	6.58	6.43	6.5	7.76	7.43	7.64	8.87	8.19	8.98	9.38	10.5	10.14	5.14	5.77	5.95	6.41	6.87	7.39	2829
THETA _{asm}	3.64	4.61	5.85	6.75	7.29	8.05	7.06	7.18	7.14	7.61	7.71	7.83	9.01	8.21	9.63	10.74	10.09	10.91	5.21	6.03	6.3	6.73	7.17	7.74	2829
THETA	3.23	3.95	4.93	5.9	6.59	6.92	6.21	6.21	6.57	6.67	6.35	6.86	7.96	7.72	8.32	9.5	8.81	9.77	4.5	5.25	5.49	5.87	6.29	6.8	2829
RBF	4.28	4.86	5.59	6.51	7.11	7.44	6.98	6.59	6.79	6.77	7.09	6.9	8.17	7.89	8.73	9.44	9.01	9.39	5.31	5.96	6.17	6.43	6.83	7.24	2829
ForcX	3.33	4.08	4.97	6.2	6.47	7.0	6.28	6.52	6.31	6.91	6.23	7.43	8.22	8.61	9.24	10.31	9.82	10.14	4.64	5.34	5.61	5.98	6.52	7.11	2829
Chronos	3.22	4.0	5.23	6.51	6.78	7.5	6.38	6.56	6.59	6.6	6.69	7.52	8.03	8.7	9.12	8.72	10.12	10.14	4.74	5.54	5.77	6.07	6.47	6.95	2829
TimeGPT	4.25	5.15	6.35	7.63	7.97	8.55	7.15	6.9	6.33	6.79	6.6	6.93	8.31	8.15	8.51	10.35	9.74	10.44	5.84	6.65	6.74	7.72	7.04	7.56	2829
Moirai	5.46	6.08	6.69	7.83	7.87	8.4	7.3	7.34	7.49	8.23	7.96	7.96	9.46	8.88	9.97	10.46	9.91	10.99	6.52	7.06	7.12	7.38	7.79	8.24	2829

Table 13: Median Relative Absolute Error with All Data.

Method	Forecast Horizon																		Average of Forecast Horizon					# obs	
	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	Avg 1-4	Avg 1-6	Avg 1-8	Avg 1-12	Avg 1-15		Avg 1-18
NAIVE2	0.18	0.2	0.25	0.31	0.32	0.35	0.32	0.33	0.32	0.33	0.3	0.29	0.33	0.32	0.34	0.35	0.36	0.37	0.24	0.27	0.28	0.29	0.3	0.31	2829
HOLT	0.16	0.19	0.23	0.28	0.29	0.31	0.29	0.28	0.24	0.3	0.24	0.25	0.27	0.26	0.29	0.31	0.32	0.33	0.22	0.24	0.25	0.26	0.26	0.27	2829
DAMPEN	0.16	0.18	0.22	0.27	0.29	0.3	0.27	0.29	0.25	0.28	0.27	0.27	0.3	0.28	0.31	0.34	0.33	0.35	0.21	0.24	0.25	0.25	0.26	0.28	2829
WINTER	0.16	0.19	0.24	0.28	0.3	0.31	0.28	0.28	0.25	0.3	0.26	0.24	0.27	0.26	0.3	0.31	0.33	0.32	0.22	0.25	0.26	0.26	0.26	0.27	2829
COMB S-H-D	0.16	0.19	0.23	0.27	0.28	0.3	0.27	0.28	0.25	0.3	0.25	0.25	0.28	0.27	0.3	0.32	0.32	0.33	0.21	0.24	0.25	0.25	0.26	0.27	2829
B-J auto	0.16	0.19	0.23	0.27	0.29	0.3	0.28	0.29	0.27	0.28	0.28	0.27	0.3	0.31	0.31	0.36	0.32	0.35	0.21	0.24	0.25	0.26	0.27	0.28	2829
AutoBox1	0.17	0.2	0.24	0.28	0.3	0.32	0.27	0.29	0.26	0.28	0.26	0.26	0.3	0.3	0.31	0.33	0.34	0.34	0.22	0.25	0.26	0.26	0.27	0.28	2829
AutoBox2	0.17	0.19	0.23	0.28	0.28	0.31	0.29	0.29	0.26	0.28	0.26	0.26	0.29	0.29	0.31	0.34	0.34	0.35	0.22	0.24	0.26	0.26	0.27	0.28	2829
AutoBox3	0.16	0.21	0.24	0.29	0.3	0.32	0.29	0.29	0.26	0.29	0.25	0.25	0.29	0.28	0.3	0.33	0.35	0.35	0.22	0.25	0.26	0.26	0.27	0.28	2829
ROBUST-Trend	0.17	0.2	0.24	0.27	0.29	0.3	0.29	0.31	0.31	0.32	0.3	0.27	0.3	0.33	0.34	0.34	0.34	0.37	0.22	0.24	0.26	0.27	0.28	0.29	2829
ARARMA	0.16	0.19	0.24	0.27	0.28	0.31	0.28	0.3	0.27	0.29	0.26	0.26	0.28	0.28	0.3	0.33	0.33	0.35	0.22	0.24	0.25	0.26	0.26	0.28	2829
Auto-ANN	0.16	0.2	0.22	0.28	0.29	0.33	0.28	0.29	0.26	0.3	0.27	0.26	0.3	0.3	0.31	0.33	0.33	0.34	0.22	0.25	0.26	0.26	0.27	0.28	2829
Flors-Pearc1	0.16	0.19	0.23	0.28	0.3	0.31	0.29	0.3	0.26	0.29	0.27	0.27	0.31	0.31	0.33	0.35	0.35	0.37	0.22	0.24	0.26	0.26	0.27	0.29	2829
Flors-Pearc2	0.19	0.22	0.25	0.29	0.3	0.31	0.28	0.29	0.25	0.28	0.26	0.29	0.3	0.3	0.32	0.33	0.34	0.36	0.24	0.26	0.27	0.27	0.28	0.29	2829
PP-Autocast	0.16	0.19	0.23	0.27	0.29	0.31	0.28	0.29	0.27	0.31	0.28	0.27	0.3	0.29	0.31	0.34	0.34	0.36	0.21	0.24	0.25	0.26	0.27	0.28	2829
ForecastPro	0.15	0.18	0.21	0.25	0.27	0.29	0.27	0.27	0.24	0.26	0.24	0.23	0.28	0.26	0.28	0.33	0.31	0.33	0.2	0.22	0.24	0.24	0.25	0.26	2829
SMARTFCS	0.17	0.2	0.24	0.28	0.29	0.31	0.28	0.28	0.25	0.32	0.27	0.27	0.3	0.27	0.31	0.32	0.34	0.34	0.22	0.25	0.26	0.26	0.27	0.28	2829
THETA _{sm}	0.17	0.2	0.24	0.28	0.3	0.32	0.28	0.3	0.28	0.31	0.28	0.27	0.3	0.27	0.3	0.32	0.32	0.32	0.22	0.25	0.26	0.27	0.27	0.28	2829
THETA	0.15	0.18	0.21	0.25	0.27	0.28	0.26	0.27	0.26	0.28	0.25	0.24	0.27	0.26	0.29	0.31	0.31	0.33	0.2	0.22	0.23	0.24	0.25	0.26	2829
RBF	0.2	0.21	0.24	0.27	0.28	0.3	0.28	0.28	0.27	0.31	0.25	0.25	0.27	0.25	0.29	0.31	0.31	0.31	0.23	0.25	0.26	0.26	0.26	0.27	2829
ForcX	0.15	0.18	0.23	0.26	0.28	0.29	0.27	0.28	0.25	0.28	0.25	0.26	0.3	0.29	0.31	0.34	0.32	0.33	0.2	0.23	0.24	0.25	0.26	0.27	2829
Chronos	0.16	0.18	0.22	0.26	0.28	0.31	0.26	0.28	0.25	0.25	0.24	0.25	0.26	0.27	0.28	0.32	0.29	0.31	0.2	0.24	0.24	0.25	0.25	0.26	2829
TimeGPT	0.2	0.21	0.27	0.3	0.33	0.36	0.28	0.3	0.25	0.28	0.25	0.26	0.31	0.28	0.3	0.34	0.33	0.36	0.24	0.28	0.28	0.27	0.28	0.29	2829
Moirai	0.24	0.26	0.28	0.3	0.32	0.35	0.31	0.32	0.28	0.34	0.29	0.28	0.33	0.31	0.33	0.35	0.34	0.37	0.27	0.29	0.3	0.3	0.3	0.31	2829

Note: Top Three Scores in Each Column are Bolded.

Appendix D

Fine-Tuning TimeGPT Experiment

TimeGPT allows for straightforward fine-tuning, offering a choice between various loss functions such as MAE, MSE, RMSE, MAPE, and SMAPE. During fine-tuning, you can specify the desired loss function and the number of iterations, which indicates how many times the model will train on the provided data before generating forecasts.

Although our thesis primarily focuses on zero-shot forecasting, we were curious about the fine-tuning capabilities of the foundation model. To explore this, we conducted a small experiment to assess its impact. This experiment was limited to monthly data, as previous evaluations indicated that the model performs poorly with yearly and quarterly data. The results showed a slight improvement, with the MASE decreasing from 0.902 to 0.894. In this experiment, we performed 10 fine-tuning iterations and did not specify a loss function, aligning with our philosophy that foundation models should be easy to use.




Figure 22: Fine-Tuned TimeGPT on Monthly Data

Appendix E

Robustness Checks

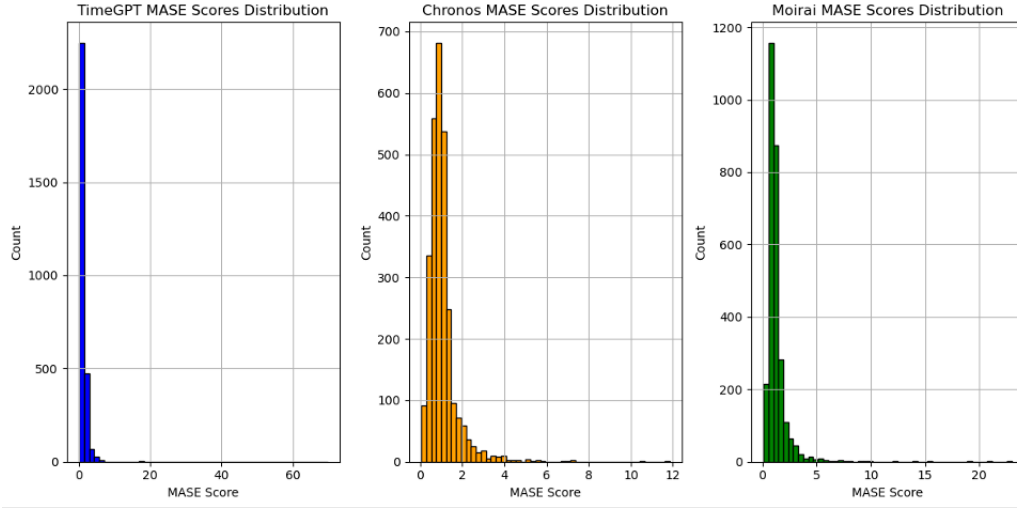


Figure 23: Distributions of MASE Scores for Foundation Models

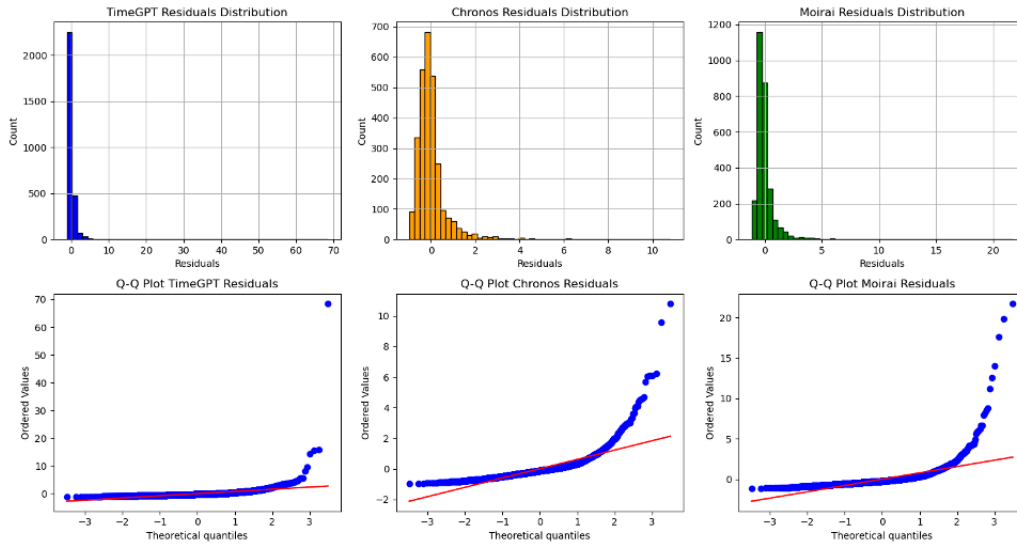


Figure 24: Distributions of Residuals for Foundation Models (MASE Scores)

Results of Friedman’s Test For the group consisting of MASE Scores for TimeGPT, Moirai and Chronos:

Friedman’s Test Statistic: 247.96818663839076, p -value: $1.426896352175751e - 54$

Conclusion: There is a statistically significant difference between the models.

Results of Levene's Test For the group consisting of MASE Scores for TimeGPT, Moirai and Chronos:

Levene's Test Statistic: 4.808090348578604, p -value: 0.008185691864333012

Reject the null hypothesis - variances are not equal (*heteroscedasticity*).

Results of Wilcoxon Signed-Rank Test Only the tests where we accept the null hypothesis when comparing a selection of models. We refer to the [GitHub](#) to see the test statistics for all comparisons drawn from the selection. The selection consists of: Theta, Naive2, TimeGPT, Chronos, Moirai, ForecastPro, ForcX, Dampen, COMB S-H-D and RBF.

Wilcoxon Signed-Rank Test between Chronos and RBF:

Test Statistic: 1962836.0, p -value: 0.3733125068089125

Fail to reject the null hypothesis - the median of the differences in the matched pairs is equal to 0

Wilcoxon Signed-Rank Test between ForecastPro and ForcX:

Test Statistic: 1770965.0, p -value: 0.14662673332067305

Fail to reject the null hypothesis - the median of the differences in the matched pairs is equal to 0

Wilcoxon Signed-Rank Test between ForecastPro and COMB S-H-D:

Test Statistic: 1969052.0, p -value: 0.4549310940170622

Fail to reject the null hypothesis - the median of the differences in the matched pairs is equal to 0

Wilcoxon Signed-Rank Test between ForcX and Dampen:

Test Statistic: 1911761.0, p -value: 0.06057003599052172

Fail to reject the null hypothesis - the median of the differences in the matched pairs is equal to 0

Wilcoxon Signed-Rank Test between ForcX and RBF:

Test Statistic: 1919268.0, p -value: 0.058352805926870525

Fail to reject the null hypothesis - the median of the differences in the matched pairs is equal to 0

Wilcoxon Signed-Rank Test between Dampen and RBF:

Test Statistic: 1961356.0, p -value: 0.3553046027801884

Fail to reject the null hypothesis - the median of the differences in the matched pairs is equal to 0

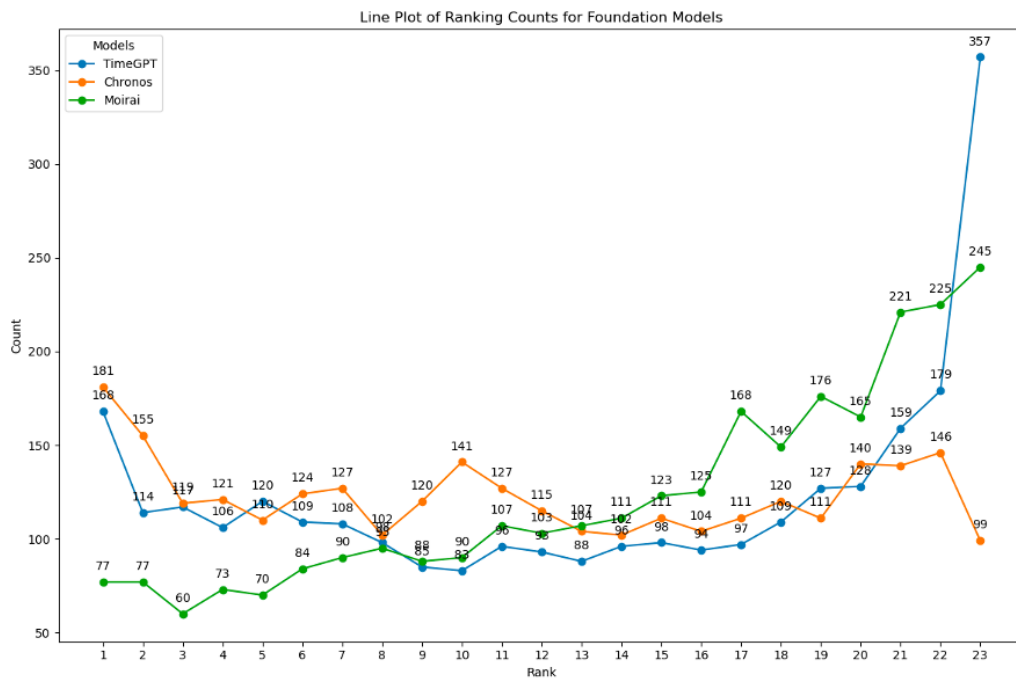


Figure 25: Counts of Ranking for Foundation Models as Line Plot




Figure 26: Heatmap of Rankings in % for Selection of Models